

Ops C00 009 - Written Answer

Thesis

Single Grain does not need more AI output. It needs an operating loop that turns AI output into accepted work: source check, score-to-route gate, owner execution, outcome log, and memory update. My first job as GM is to protect revenue, reduce Eric's review load, and make Single Brain compound instead of creating review debt.

Part 1 - First 48 Hours

I would create one command queue across Mission Control, Slack, CRM, ClickFlow, Oracle outputs, and expired outreach. I would not start with stakeholder interviews.

Priority order: ClickFlow churn first, because 2 lost customers this month `[Observed from brief]` is active retention damage; stale prospects second, because 3 prospects went cold `[Observed from brief]`; queue routing third, because 23 pending items `[Observed from brief]`, some 4 days old `[Observed from brief]`, need sorting but age alone is not priority; Oracle quick wins fourth, because 15 keyword recommendations `[Observed from brief]` need source checks before assignment.

First day: freeze new auto-approvals except reversible internal summaries, snapshot the queue, and route every item as `dead`, `needs\_source\_check`, `green\_auto`, `yellow\_owner`, `yellow\_council`, or `red\_eric`. Churn gets usage, last touch, ticket history, contract value, and save path. Stale prospects get stage, source, last reply, owner, and re-entry note or dead mark. By end of day, every item has an owner or rejection reason.

Eric's 30 minutes `[Observed from brief]`: confirm fire order, name red lines, define what I can decide for 7 days `[Assumed]`, review only 3-5 red decisions `[Observed from brief]`, and pick tomorrow's CEO leverage metric.

Part 2 - Risk Tiering

Every output gets two passes `[Assumed]`: action value and risk. Action value scores revenue proximity, time sensitivity, source confidence, owner clarity, effort size, and learning value on a 0/1/2 scale `[Assumed]`. Risk flags are irreversible customer commitment, strategic relationship, public/brand risk, legal/privacy/security, financial exposure, novelty, and low source confidence.

Green: value ≥ 7 `[Assumed]`, source verified, reversible, known pattern, no red flags. Examples: verified existing-page SEO fixes, approved-template Flash repurposing, approved non-strategic Cyborg follow-up, and Alfred internal briefings.

Yellow: value ≥ 5 `[Assumed]` or uncertainty without CEO-specific risk. Owner review handles domain judgment; expert councils handle high-value ambiguity like creative refreshes, channel pivots, SEO positioning, or paid media calls.

Red: Eric only. Pricing, refunds, contracts, strategic accounts, public claims, legal/privacy/security, senior people decisions, new core crons, killing core crons, or CEO-level priority tradeoffs. This moves Eric from 15+ daily review items `[Observed from brief]` toward 3-5 real decisions `[Observed from brief]`.

When green fails: reverse if possible, notify the owner, move that item type to yellow, add a memory record, and change the rule.

Part 3 - Cron Audit

At 30 days `[Observed from brief]`, I would audit the 48 crons `[Observed from brief]` by accepted work, not activity: action rate, revenue proximity, review

burden, source quality, owner clarity, false-positive cost, false-negative cost, and memory feedback `[Assumed]`.

Kill/pause: Competitor Pricing Monitor at 3% action `[Observed from brief]`; daily pricing noise creates executive distraction. Candidate Outreach Follow-up at 5% response `[Observed from brief]`; automated candidate nagging creates brand risk and should become weekly top-5 human review `[Assumed]`.

Improve urgently: Churn Signal Detector at 20% investigated `[Observed from brief]`; false negatives are expensive and the response loop is broken. Flags become same-day `[Assumed]` yellow escalations with owner, source links, customer state, and save-play status.

Improve/tighten: GSC Quick Win Scan at 12% actioned `[Observed from brief]` surfaces only 3-5 verified existing-page opportunities `[Assumed]`. LinkedIn Prospect Sourcer at 8% meeting conversion `[Observed from brief]` tightens ICP/trigger filters and batches daily `[Assumed]`. Keep Deal Revival at 60% followed up `[Observed from brief]`; keep/improve Content Repurpose at 35% published `[Observed from brief]` and Weekly Content Calendar at 50% used `[Observed from brief]`, but require outcome logging.

Part 4 - Operating Edge

I built an eval loop for a HeavisideOS Slack assistant that answers client-status, service, and root-cause questions from customer context. Baseline benchmark: 39 active mapped customers `[Observed internal eval]`, 117 cases `[Observed internal eval]`, overall 4.055/5.000 `[Observed internal eval]`, retrieval 3.255/5.000 `[Observed internal eval]`, gate `review\_required` `[Observed internal eval]`. After service-state, tenant communication retrieval, and Slack answer improvements, the gate passed, retrieval reached 4.964/5.000 `[Observed internal eval]`, and all 117 retrieval/citation checks passed `[Observed internal eval]`. A later structured-analysis pass reached 4.801/5.000 overall `[Observed internal eval]`.

The lesson: aggregate AI quality can hide a binding category failure. You need category gates, source checks, repeatable reports, and memory updates. Limit: HeavisideOS is not yet a fully mature company OS; this proves measured assistant improvement and retrieval/action-loop discipline.

For Oracle's 12 quick-win keywords `[Observed from brief]`, I would verify source data, existing page fit, commercial intent, effort, risk, owner, and feedback path. Existing-page changes with source proof go green; new pages, mixed intent, or strategic positioning go yellow; public claims or customer-sensitive pages go red.

In 2028 `[Observed from brief]`, the GM owns the operating loop around agents: what executes alone, what needs review, what stays human, how accepted work is measured, how failed outputs become better rules, and how the CEO's attention stays reserved for trust, pricing, strategic accounts, people, and public judgment.

Operational Experience

My operating pattern is turning ambiguous work into gates. In the M365 -> Google Workspace migration, work spanned mail, calendar, DNS, aliases, OneDrive/SharePoint, Missive, GHL, phone routing, Windows identity, and Microsoft cleanup. I installed one control center with lane state, accepted/gated status, rollback rules, and a hard rule that cleanup could not run until dependency gates passed. In PVM/GDM paid media, I used the same rule: direct GHL attribution and direct platform reads outranked stale dashboard numbers, so budgets and cuts followed source truth rather than convenient reporting. This prevented cleanup from outrunning dependency proof and kept budget/scaling decisions tied to source-truth performance data instead of stale dashboards.

MMPI-Style Self-Assessment

I am high-drive, systems-first, and comfortable with ambiguity. The strength is that I turn messy work into gates, owners, source records, and repeatable loops. The risk is overbuilding or pushing process faster than a team can adopt it. I compensate with lightweight routing: green executes, yellow gets owner/council review, red goes to the CEO.