

# ai-usage-disclosure

---

# Ops COO 009 - AI Usage Disclosure v2

## ## Tools Used

- Codex / ChatGPT-style assistant for research organization, critique, drafting, compression, and QA.
- Local repo inspection for the challenge brief, `SCORING.md`, `ops-coo-009/scoring\_rubric.md`, and `ops-coo-009/claude\_baseline.md`.
- SloaneVault/qmd search for internal source records, project notes, Eric/Single Grain research, x-accel references, HeavisideOS references, and proof candidates.
- Web research for official Single Grain GM role facts and source links.
- x-accel authenticated X capture stored in SloaneVault for Eric's public commentary, refreshed on 2026-05-13.
- Hermes/open-agents scheduler-state review for bounded cron-fleet proof.
- Decision-council review and peer-review synthesis to stress-test the packet before final packaging.
- Local PDF export / QA scripts for page-fit, number-label review, stale-claim scan, and sensitive-info scan.

## ## What AI Helped With

- Comparing the draft against Claude's published baseline.
- Organizing the answer around the GM job mandate, not just the challenge mechanics.
- Drafting and compressing the written answer, operating artifact, evidence log, and disclosure.
- Stress-checking source labels, proof tiers, failure modes, human boundaries, and score-losing claims.
- Separating HeavisideOS, x-accel, SloaneVault, and decision councils accurately.
- Identifying the need to add bounded Hermes scheduler-fleet evidence without overclaiming cron quality.

## ## What I Personally Decided

- The core thesis: AI output only matters when it becomes accepted work, owner action, memory, margin, retention, and revenue/FTE leverage.
- The first 48-hour priority order: ClickFlow churn first, stale prospects second, queue routing third, Oracle keyword source checks fourth.
- The green/yellow/red routing model: green only for reversible, source-verified, known-pattern work; yellow for owner/council judgment; red for Eric-level trust, pricing, strategy, public, legal, people, and key relationship decisions.
- The cron audit posture: kill or pause low-action/high-noise workflows instead of keeping every automation alive.
- The evidence posture: lead with Heaviside RPE/profitability proof, support with HG Market Report, HG SEO Commander, SloaneVault, x-accel, decision councils, and carefully bounded HeavisideOS proof.
- The cron proof posture: use Hermes/open-agents as evidence of operating a monitored scheduler fleet, not as proof that every scheduled job deserves to exist.
- The maturity boundary: x-accel is separate from HeavisideOS; HeavisideOS is early practical Slack/client/performance summarization plus company-brain direction, not a fully adopted company operating system.
- The redaction rule: raw internal reports, P&L exports, CRM data, Slack messages, and private vault notes should not be attached externally without redaction.

## ## What I Checked

- `challenges/ops-coo-009/brief.md`
- `challenges/ops-coo-009/scoring\_rubric.md`
- `SCORING.md`
- `challenges/ops-coo-009/claude\_baseline.md`
- GM role posting facts for compensation and remoteness.
- Heaviside RPE/profitability calculations from QuickBooks P&L snapshots and labor-to-FTE method.
- Applicant interview notes for HG Market Report, HG SEO Commander, operating style, self-assessment, and approved wording.
- SloaneVault project notes and lessons for corrections and maturity boundaries.
- Number labels in the v2 written answer and artifact.
- Sensitive-info risk in internal proof artifacts.
- Hermes scheduler manifests and live-state counts used in the redacted cron-fleet summary.

## ## Known Limits

- GM role facts should be rechecked immediately before upload because public job pages can change.
- Eric/X commentary is based on captured corpora refreshed through 2026-05-13 17:52 UTC, not a guarantee of later posts.
- Heaviside RPE/FTE history uses observed P&L revenue and net income, with FTE estimated from labor spend using documented cost/FTE benchmarks.
- HG Market Report and HG SEO Commander time savings are applicant estimates unless later backed by time logs.
- Hermes/open-agents scheduler proof is internal source-record evidence and does not independently verify business value for every scheduled job.
- Internal proof summaries should be attached only in redacted form.