

evidence-log

Ops COO 009 - Evidence Log v2

Proof tier reference from `SCORING.md`:

- Tier 0: claims only.
- Tier 1: screenshots.
- Tier 2: demo artifact.
- Tier 3: logs or source records.
- Tier 4: before/after data.
- Tier 5: independent verification.

This evidence log maps the v2 written answer and operating artifact to proof tiers and source labels.

Evidence Table

[Claim | Number source labels | Proof tier | Evidence / source | Shareability / caveat |
[---][---][---][---]
The `ops-coo-009` brief describes `48` crons, `50+` daily messages, `15+` daily review items, `23` pending items, `2` ClickFlow churn customers, `3` stale prospects, `15` inherited Oracle quick-win keywords, and a separate Part 4 evaluation prompt for `12` quick-win keywords from last night.	[Observed from brief]	3	`challenges/ops-coo-009/brief.md`	Shareable; public repo brief.
The `ops-coo-009` brief says Single Grain's TES/SaaS products are targeting `50%` TES revenue by 2028.	[Observed from brief]	3	`challenges/ops-coo-009/brief.md`	This is the strategic reason revenue/FTE should rise through productized/tech-enabled revenue and accepted agent work, not artificial shrinkage.
Eric's intended load is `3-5` decisions/day and `15` minutes/day, not `15+` review items/day or `60+` minutes/day.	[Observed from brief]	3	`challenges/ops-coo-009/brief.md`	Shareable; public repo brief.
The brief's CEO profile says Eric is ENTJ / `8w7`, Kolbe Quick Start `9`, Follow Thru `2`, self-rated weak on follow-through, organization, and workflow consistency, and starts `10` things while finishing `3`.	[Observed from brief]	3	`challenges/ops-coo-009/brief.md`	Treat as the central operating clue, not trivia.
The GM role is remote and compensation is `\$185K-\$225K DOE`.	[Observed from GM posting]	3	`https://careers.singlegrain.com/jobs/567472-general-manager-agency`, verified 2026-05-14.	Recheck before upload because public job pages can change.
Single Grain's GM mandate is to increase revenue per employee from `\$180K` to `\$300K-\$400K`.	[Observed from GM posting]	3	`https://careers.singlegrain.com/jobs/567472-general-manager-agency`, verified 2026-05-14.	Shareable; public role posting.
Single Grain publicly positions the agency around AI-native execution, pipeline/revenue accountability, and scaling results without scaling headcount.	[Observed from public websites]	3	`https://www.singlegrain.com/` and `https://www.singlegrain.com/landing/`, verified 2026-05-13.	Use as positioning/context, not as internal performance proof.
Single Brain is publicly described as a managed AI agent team inside Salesforce, Gong, GA4, HubSpot, and Slack.	[Observed from public websites]	3	`https://www.singlegrain.com/ai-tools/`, verified 2026-05-13.	Public product/positioning page.
The standalone Single Brain site frames the architecture as one trusted brain, one commander, and a specialist fleet with isolated client deployments and an SSR trust standard.	[Observed from public websites]	3	`https://singlebrain.com/`, verified 2026-05-13.	Public product/positioning page; do not claim actual internal deployment details beyond the page.
ClickFlow publicly supports a strategy -> writing -> reporting loop and developer-accessible workflows driven by GSC/GA4, CMS publishing, MCP/REST, and AI visibility data.	[Observed from public websites]	3	`https://www.clickflow.com/`, `https://docs.clickflow.com/`, and `https://docs.clickflow.com/capabilities/analytics`, verified 2026-05-13.	Public docs/product pages.
Karrot publicly supports LinkedIn ABM personalization, HubSpot pipeline reporting, deal activity timelines, and personalized landing pages.	[Observed from public websites]	3	`https://karrot.ai/` and `https://www.singlegrain.com/ai-marketing-tools/`, verified 2026-05-13.	Public product pages.
Single Grain's AI Revenue Agents guide emphasizes data quality, bidirectional CRM/marketing automation sync, staged rollout, approval rules, escalation thresholds, recovery protocols, and KPIs.	[Observed from public websites]	3	`https://www.singlegrain.com/artificial-intelligence/ai-revenue-agents/`, verified 2026-05-13.	Public thought-leadership page; use for calibration.
Single Grain's Company Intelligence guide frames the system as an active intelligence layer across customer behavior, internal ops data, predictive modeling, multi-agent execution, governance, and audit logging.	[Observed from public websites]	3	`https://www.singlegrain.com/artificial-intelligence/company-intelligence/`, verified 2026-05-13.	Public thought-leadership page; use for calibration.
Eric publicly wrote that Single Grain has rebuilt operations around Claude Code/OpenClaw and saves `90+` hours/month on workflows that previously required dedicated headcount.	[Observed from public commentary]	3	`https://newsletter.levelingup.com/p/how-openclaw-claude-code-will-shift-your-business-completely`, verified 2026-05-13.	Public founder commentary, not a verified role fact.
Heavyside started in 2011 and Chris remains founder/operator.	[Observed applicant claim]	0-1	Applicant interview notes: `tasks/ops-coo-009-applicant-interview-notes.md`.	Applicant claim; confirm before final upload.
Chris spent the last `2` years re-engineering Heavyside, PVM, and GDM around an intelligence layer, software-driven workflows, and CLI-driven operations.	[Observed applicant claim]	0-2	Applicant interview notes; supporting repo/vault artifacts can be summarized if approved.	Applicant claim; share only high-level descriptions unless artifacts are redacted.
Chris sustained and grew Heavyside while holding a full-time PASCO role owning SEO/digital marketing and now ML/AI work.	[Observed applicant claim]	0-2	Applicant correction on 2026-05-14; candidate profile and written answer.	Strong talent-leadership and delegation claim. Frame as delegated operating-capacity proof, not divided-attention proof. Applicant-stated unless backed by calendar, PASCO role evidence, operator references, or org chart.
Chris's existing businesses take less than `5` hours/week to manage because he hires competent operators, gives ownership, and avoids micromanagement.	[Observed applicant claim]	0-2	Applicant correction on 2026-05-14; candidate profile and written answer.	Strong talent-leadership, delegation, and new-role capacity claim. Applicant-stated unless backed by calendar/time logs, operator references, or org chart.
Heavyside has a public operating bench across technical operations, web, SEO, PPC, operations, project flow, client relationship, client success, and reporting/admin support.	[Observed public website / internal ownership map]	3	`https://heavysidegroup.com/about`, `https://heavyside.digital/about`, checked 2026-05-14; summarized in `heavyside-org-chart-redacted-summary.md`.	Supports talent leadership, delegation, leadership bench, and founder leverage. Not an FTE or payroll-validation artifact.
Heavyside moved from `\$55.6K` revenue/FTE and `2.6%` net margin in 2023 to `\$77.2K` and `14.2%` in 2024, `\$90.0K` and `12.9%` in 2025, and `\$120.1K` and `17.1%` in 2026 Q1.	[Observed internal P&L / Estimated FTE]	4	Native QuickBooks ProfitAndLoss snapshots in Agency Financials plus labor-spend-to-FTE model; summarized in `heavyside-rpe-profitability-redacted-summary.md`, with internal working calculation available for walkthrough.	Attach the redacted summary only. Do not attach raw confidential P&L exports.
Current active-plus-paused Heavyside MRR is about `\$70.7K/mo`, or `\$848K` annualized, and roughly `\$126K` revenue per paid-team FTE using `6.71` paid-team FTE.	[Observed internal customer/payroll records]	3-4	Agency Financials DB query and payroll/FTE calculation summarized in `tasks/ops-coo-009-heavyside-rpe-calculation.md`.	Useful backup; main written answer uses historical P&L basis instead.
HG Market Report implements a lead/report/presentation workflow: business and Google Business Profile capture, AI keyword/service-area setup, SEO audit / SERP / geogrid jobs, derived lost-revenue / quick-win / competitor views, public presentation URL, and report-ready email.	[Observed from private repo docs]	3	`hg-market-report` README, `docs/embed/README.md`, `docs/reports/report-architecture.md`, presentation routes, summarized in `heavyside-sales-delivery-automation-redacted-summary.md`.	Summarize only; do not attach customer reports, private admin screens, env examples, or source excerpts without review.
HG Market Report saves roughly `1-2` hours per standard sales audit.	[Estimated applicant claim]	0-2	Applicant estimate tied to replaced steps: manual source pulls, local search checks, screenshot/report assembly, presentation creation, and first delivery.	Label as estimate; stronger if later backed by time logs or before/after task records.
HG SEO Commander automates large parts of recurring SEO delivery across GA4, Search Console, Google Business Profile, SERP/geogrid tracking, citations, AI articles, GBP posts, bulk-posting approvals, scheduled syncs, BullMQ queues, and CLI/admin tools.	[Observed from private repo docs]	3	`hg-seo-commander` README, CLI docs, GBP implementation assessment, scheduled-sync notes, summarized in `heavyside-sales-delivery-automation-redacted-summary.md`.	Summarize only; strategy and quality review remain human. Stronger with redacted Loom or demo report.
HG SEO Commander saves `3-4` hours per active SEO client per month.	[Estimated applicant claim]	0-2	Applicant estimate tied to routine data pulls, status checks, report assembly, GBP/content prep, and queue recovery.	Label as estimate; stronger if later backed by time logs.
The GBP-post bug/process example shows process enforcement through CLI tooling and adoption by making the operator more capable.	[Observed applicant claim]	0-2	Applicant interview notes.	Good narrative proof; not independently verified yet.

| SloaneVault is a durable memory/source-record system used across projects. | "[Observed internal source]" | 3 | SloaneVault/qmd index, project notes, lessons files, and decision records. | Share high-level structure only, not sensitive vault contents. |

| x-accel provides a proven scoring/gate analogue for AI-generated content before action. | "[Observed internal source]" | 3 | x-accel publish-gate source records and redacted workflow summary. | x-accel is separate from HeavisideOS; summarize unless code excerpts are approved. |

| Decision councils are a reusable judgment mechanism for high-value ambiguous decisions. | "[Observed internal source]" | 2-3 | Decision Council workflow source record and safe/anonymized council artifact. | Use only as supporting proof; do not make every decision a council. |

| HeavisideOS is currently used in practice for SlackBot summaries/query flows around recent client communications and paid ads/SEO/business-performance context. | "[Observed internal source]" | 3 | HeavisideOS Slack connector source records and project lesson. | Do not claim HeavisideOS runs the company end-to-end. |

| HeavisideOS Slack assistant eval improved from overall `4.055/5.000` to `4.467/5.000`, later `4.801/5.000`, and retrieval improved from `3.255/5.000` to `4.964/5.000` across `117` cases. | "[Observed internal eval]" | 4 | Redacted HeavisideOS Slack assistant eval summary; raw report IDs `20260226-225557`, `20260227-015739`, and `20260227-025448`. | Optional proof. Submit redacted metrics only; raw reports contain customer names. |

| Chris operates a Hermes/open-agents scheduler fleet with `88` repo-checked jobs, `28` Hermes skills, and `3` managed contexts; live fleet status on 2026-05-14 reported healthy across `3/3` nodes. | "[Observed internal scheduler state]" | 3 | Internal Hermes/open-agents manifests, `sync-hermes.py --verify`, and `hermes-fleet.py status --json`; summarized in `hermes-cron-fleet-redacted-summary.md` and `hermes-closed-loop-agent-proof.md`. | Bounded cron-fleet operating evidence. `sync-hermes.py --verify` showed no job issues but did show `4` skill content mismatches; do not claim perfect fleet hygiene or every job's business value. |

| Eric/X public commentary emphasizes Single Grain/company-brain workflows, revenue agents, service-as-software, proof-of-work hiring, Slack/Teams adoption, and model/tool-spend governance. | "[Observed captured X corpus]" | 3 | Authenticated x-accel pull of `@ericosiu`, refreshed 2026-05-13 through 17:52 UTC. | Public commentary, not role fact; do not overfit current copy to stale posts. |

| The proposed score-to-route gate can reduce Eric's default review burden. | "[Assumed operating design]" | 2 | `tasks/ops-coo-009-operating-artifact-v2.md`. | Proposed artifact, not observed Single Grain result. |

| The cron audit kills/pauses `2` sample crons and improves/keeps the rest based on action rate and review debt. | "[Assumed verdict from observed brief data]" | 2-3 | `tasks/ops-coo-009-operating-artifact-v2.md`; sample cron data from public brief. | Verdicts are proposed; action rates are brief facts. |

| MMPI-style self-assessment: founder first; INTP / `6w5` with strong Type `1`; Kolbe `9-4-3-2`; high agency, detail-oriented, systems-minded, ambiguity-tolerant; risk is overbuilding/capability sprawl. | "[Observed applicant self-assessment]" | 1-2 | "[[30 - Personal/37 - Knowledge Graph/people/chris-sloane/personality-and-operating-profile]]", applicant interview notes, candidate profile addendum, and v2 written answer. | Self-assessment, not validated clinical psychometrics; applicant should approve before final upload. |

| AI usage disclosure. | "[Observed workflow]" | 1-2 | `tasks/ops-coo-009-ai-disclosure-v2.md`. | Submit final disclosure. |

Claims To Avoid Or Demote

- Do not claim HeavisideOS is a mature company-wide operating system.
- Do not claim x-accel is part of HeavisideOS.
- Do not claim Heaviside is already at Single Grain's `\$300K-\$400K` revenue/FTE target.
- Do not imply Heaviside's FTE reduction was artificial shrinkage; frame it as sensible simplification after overstaffing, supported by improved profitability.
- Do not claim direct agency-scale equivalence to Single Grain's `45`-person organization without a verified source example.
- Do not claim Hermes/open-agents proves every scheduled job is valuable; use it only as bounded evidence of operating a monitored scheduler fleet with pause/error visibility.
- Do not attach raw HeavisideOS eval reports, raw vault notes, private client data, customer screenshots, or raw P&L exports externally without redaction.

Final Recheck

- Reconfirm the GM posting before upload.
- Confirm or replace LinkedIn URL.
- Re-scan attachments for customer names, private client data, secrets, tokens, and raw report tables.
- Confirm applicant approval of self-assessment and operating-experience claims.
- Re-run number-label audit after edits.