

heavysideos-slack-eval-redacted-summary

HeavysideOS Slack Assistant Eval - Redacted Summary

This is a redacted evidence summary for `ops-coo-009`. It is designed to be shareable without raw customer names or raw report tables.

Why This Is The Primary Part 4 Proof

This is the cleanest operating-edge proof because it shows a real AI workflow with:

- a repeatable eval runner,
- production API benchmark reports,
- explicit scoring dimensions,
- gate status,
- before/after results,
- code changes tied to the improvement,
- and a documented learning loop.

It should be used as proof of measured AI assistant improvement, retrieval discipline, and operating-loop behavior. It should not be used to claim HeavysideOS is a fully mature company operating system.

Source Records

Evidence	Source reference	Shareability
Eval runbook	HeavysideOS Slack Assistant Usefulness Eval Program	Available as summarized/redacted source record.
Baseline report	`slack-usefulness/20260226-225557/report.md`	Raw report contains customer names; not attached.
Final pass report	`slack-usefulness/20260227-015739/report.md`	Raw report contains customer names; not attached.
Later structured-analysis pass	`slack-usefulness/20260227-025448/report.md`	Raw report contains customer names; not attached.
Source commit	HeavysideOS commit `1dd49b8`	Shareable if repo/source access is appropriate.
Eval runner	`scripts/assistant-evals/run-slack-assistant-eval.ts`	Shareable if repo/source access is appropriate.
Eval config	`config/evals/slack-assistant/usefulness-suite.json`	Shareable if tenant data is not embedded.

Redacted Metrics Table

Metric	Baseline	Final pass	Delta
Active mapped customers	39 `[Observed internal eval]`	39 `[Observed internal eval]`	no change
Questions per tenant	3 `[Observed internal eval]`	3 `[Observed internal eval]`	no change
Total eval cases	117 `[Observed internal eval]`	117 `[Observed internal eval]`	no change
Overall score	4.055 / 5.000 `[Observed internal eval]`	4.467 / 5.000 `[Observed internal eval]`	+0.412 `[Observed internal eval]`
Root-cause score	5.000 / 5.000 `[Observed internal eval]`	5.000 / 5.000 `[Observed internal eval]`	0.000 `[Observed internal eval]`
Evidence score	4.872 `[Observed internal eval]`	4.859 `[Observed internal eval]`	-0.013 `[Observed internal eval]`
Actionability score	3.705 `[Observed internal eval]`	3.991 `[Observed internal eval]`	+0.286 `[Observed internal eval]`
Retrieval score	3.255 `[Observed internal eval]`	4.964 `[Observed internal eval]`	+1.709 `[Observed internal eval]`
Formatting score	2.987 `[Observed internal eval]`	3.000 `[Observed internal eval]`	+0.013 `[Observed internal eval]`
Gate status	`review\_required` `[Observed internal eval]`	`pass` `[Observed internal eval]`	gate cleared
API responses passing	not recorded in baseline table	117 `[Observed internal eval]`	final pass complete
Retrieval non-empty pass	baseline had 3 failures `[Observed internal eval]`	117 pass / 0 fail `[Observed internal eval]`	tenant retrieval gap cleared
Citation presence pass	baseline had 3 warnings `[Observed internal eval]`	117 pass `[Observed internal eval]`	citation warnings cleared

Later structured-analysis pass (`20260227-025448`): overall score reached 4.801 / 5.000 `[Observed internal eval]`, actionability reached 5.000 / 5.000 `[Observed internal eval]`, retrieval remained 4.964 / 5.000 `[Observed internal eval]`, and the gate remained `pass` `[Observed internal eval]`.

What Changed

- Implemented a Slack assistant usefulness eval runner and config.
- Added service-focused golden cases across active PVM/GDM customers.
- Implemented baseline service architecture: catalog defaults, tenant overrides, and service fact overrides.
- Added Slack `SERVICE` fact-update path with tenant service sync on approval.
- Remediated tenant-specific communications retrieval gaps with targeted ClickUp/Missive sync.
- Added a safe Slack test lane so live Slack-path validation can run without posting into customer channels by default.

Submission Use

Use this as the primary answer to: "What's one operational workflow you've personally automated or augmented with AI? Give us the before/after with real metrics."

Recommended phrasing:

> I built an eval loop for a Slack assistant that answers client-status, service, and root-cause questions from customer context. The first production benchmark across 39 active mapped customers `[Observed internal eval]` and 117 cases `[Observed internal eval]` showed an overall score of 4.055/5.000 `[Observed internal eval]` and retrieval of 3.255/5.000 `[Observed internal eval]`, with the gate requiring review. After targeted retrieval/service-state fixes and tenant sync remediation, the full pass reached 4.467/5.000 overall `[Observed internal eval]`, 4.964/5.000 retrieval `[Observed internal eval]`, and a passing gate `[Observed internal eval]`. The lesson was that aggregate quality can hide a binding category failure, so I added category gates, source checks, and repeatable regression reports.

Caveats

- Raw reports include customer names; external packet should use this redacted metrics table.
- The workflow proves measured assistant improvement and retrieval/action-loop discipline.
- It does not prove HeavysideOS is fully mature as a company-wide operating system.
- It maps to Single Grain because the problem is structurally similar: Slack/data-connected AI outputs must be measured by usefulness, retrieval, accepted work, and closed-loop improvement.