

score-review

Ops COO 009 - Packet v2 Score Review

Status: score review after rebuilding the v2 packet, adding public-website calibration, adding the Eric/Chris operating-profile match, adding redacted proof summaries, completing the council remediation passes, adding explicit talent-leadership / accountability proof, adding the PASCO/full-time-role delegation proof, adding the less-than-`5`-hours/week founder-management claim, and adding line-by-line job-post/challenge-compliance scoring.

Last rescored: 2026-05-14.

Reviewed files:

- `tasks/final-upload-v2/written-answer.md`
- `tasks/final-upload-v2/written-answer.pdf`
- `tasks/final-upload-v2/operating-artifact.md`
- `tasks/final-upload-v2/evidence-log.md`
- `tasks/final-upload-v2/ai-usage-disclosure.md`
- `tasks/final-upload-v2/candidate-profile.md`
- `tasks/final-upload-v2/attachment-map.md`
- `tasks/final-upload-v2/job-post-line-scorecard.md`
- `tasks/final-upload-v2/challenge-compliance-audit.md`
- `tasks/final-upload-v2/application-answers.md`
- `tasks/final-upload-v2/cover-letter.md`
- `tasks/final-upload-v2/hermes-closed-loop-agent-proof.md`
- `tasks/final-upload-v2/heaviside-org-chart-redacted-summary.md`
- `tasks/final-upload-v2/heaviside-rpe-profitability-redacted-summary.md`
- `tasks/final-upload-v2/heaviside-sales-delivery-automation-redacted-summary.md`
- `tasks/final-upload-v2/heavisideos-slack-eval-redacted-summary.md`
- `tasks/final-upload-v2/hermes-cron-fleet-redacted-summary.md`

Scoring sources:

- Live GM job post: `https://careers.singlegrain.com/jobs/567472-general-manager-agency`, rechecked 2026-05-14.
- `challenges/ops-coo-009/claude\_baseline.md`
- `challenges/ops-coo-009/scoring\_rubric.md`
- `SCORING.md`

Mechanical QA

Check	Result
Written answer word count	`1724` words
Written answer compact PDF	`tasks/final-upload-v2/written-answer.pdf`
PDF page count	`2` pages
QA helper	Pass
Candidate/blocker language	Pass
Required sections	Pass: What stays human, What breaks, AI usage, Evidence, MMPI, LinkedIn
Local path scan	Pass
Sensitive-info scan	Warning only: redaction warnings in evidence/proof files; no secrets or raw client data detected by helper.

Score Posture

Rubric gate	Read
Real judgment	Strong. Clear triage order, hard cron kills, CEO decision boundary, and explicit tradeoffs.
Execution artifact	Strong. Artifact includes command queue, gate, cron audit, dashboard, memory schema, cadence, human boundary, and failure modes.
Source-labeled specificity	Strong. Written answer and artifact label substantive numbers. Final human review should still scan after edits.
Proof quality	Strong. Evidence log now points to a Tier 4 RPE/profitability artifact, a Tier 4 HeavisideOS eval artifact, Tier 2-3 sales/delivery automation artifacts, Tier 3 Hermes scheduler-fleet evidence, and Tier 3 public-roster / ownership-map talent proof.
AI fluency with limits	Strong. Shows automation, source checks, routing, memory, councils, and human boundaries without overclaiming HeavisideOS maturity.
Failure handling	Strong. Written answer and artifact both name failure modes and response mechanisms.
Communication	Strong. Written answer is direct and fits `2` pages. Artifact is longer but inspectable.
Public-site fit	Strong. Recommendations now reflect Single Grain's public model: AI execution inside Slack/CRM/analytics/CMS, Single Brain command/fleet architecture, ClickFlow roadmap workflows, Karrot/HubSpot pipeline writeback, and SSR trust standards.
CEO operating fit	Strong. Packet now addresses the brief's explicit Eric profile: high Quick Start, low Follow Thru, weak workflow consistency, `10` starts / `3` finishes, and the need for a GM who turns starts into closed loops.
Talent leadership / accountability	Stronger. Written answer and candidate profile now state the founder-leverage proof directly: Chris sustained and grew the agency while holding a full-time PASCO role, and his existing businesses take less than `5` hours/week to manage because he hires competent operators, gives ownership, avoids micromanagement, and intervenes where ownership or performance requires it. The safer frame is delegated operating capacity, not divided attention.
TES strategy	Stronger. Written answer now names the `50%` TES revenue by 2028 goal and connects revenue/FTE to productized/tech-enabled revenue, accepted agent work, and disciplined core services.
Complexity reduction	Good but watch. The thesis and artifact reduce review debt, but the support stack can look complex. Use the lean upload hierarchy and keep Hermes/x-accel/councils as supporting proof unless requested.

GM Job Requirement Alignment

Live job-post facts rechecked 2026-05-14: Single Grain wants a high-bar GM who builds operational leverage, moves revenue per employee from `\$180K` to `\$300K-\$400K`, reduces CEO dependency, operates the agency, enforces accountability, manages profitability and hiring discipline, builds leadership teams, and creates standardized processes, KPI dashboards, performance-management systems, and AI-enabled workflows.

Job-post requirement	Fit score	Read
Builder of operational leverage, not traditional agency operator	`9/10`	Strong. Packet leads with RPE/margin, operating systems, command queue, gates, memory, and automation-to-accepted-work.
Increase revenue per employee while maintaining business health	`8.75/10`	Strong but not complete proof at Single Grain scale. Heaviside shows RPE and margin improvement together; target operating model is clear. Needs live SG economics in first 30 days.
Reduce CEO operational dependency	`9.5/10`	One of the strongest areas. The whole artifact is designed around reducing Eric to red decisions and CEO-level judgment.

Operated or scaled a performance marketing agency / service business	`8.5/10`	Strong founder/operator proof since 2011, but smaller than Single Grain and partially applicant-stated.
Built operational systems that improved productivity	`9.25/10`	Strong. HG Market Report, HG SEO Commander, HeavisideOS eval, Hermes scheduler, and RPE/profitability proof reinforce each other.
Demonstrated financial discipline	`8.75/10`	Strong. P&L-based RPE/margin history is exactly job-relevant. Could be even stronger with hiring-decision examples and budget governance examples.
Built leadership teams capable of operating independently	`8.75/10`	Improved. Public roster, ownership map, org chart, PASCO/full-time-role proof, and less-than-`5`-hours/week management claim now support this. Remaining weakness is lack of third-party/operator references.
Talent leadership and accountability	`8.75/10`	Stronger than before. The hiring/delegation model is clear; still could use one concrete underperformance/coaching example.
Systems and AI infrastructure	`9.5/10`	Very strong. This is probably the clearest match to the role's "tech-enabled services" direction.
Ability to take the role without dropping existing businesses	`8.75/10`	Stronger with the less-than-`5`-hours/week management claim. Keep the emphasis on delegated operating capacity and explicit boundaries, not casual split attention.

Estimated job-post fit after the latest changes: `8.85-9.15/10`.

Detailed line-by-line job-post scoring:

- `job-post-line-scorecard.md`

Detailed challenge-part and win/lose audit:

- `challenge-compliance-audit.md`

Internal Score Read

Not an official point score, but the current packet posture against the public rubric:

Rubric area	Current read
Real judgment	`9/10`
Execution artifact	`9/10`
Source-labeled specificity	`9/10`
Proof quality	`9/10`
AI fluency with limits	`9-9.25/10`
Failure handling	`8.75-9/10`
Communication	`8.25-8.75/10`
Job-post fit	`8.85-9.15/10`
Baseline-beating likelihood	High, assuming reviewer accepts the redacted evidence summaries and attachments as credible and the existing-businesses point is framed as delegated operating capacity.

Second Council Packaging Review

The second council did not recommend a rewrite. It converged on a packaging decision:

- Freeze the written answer except typo, sensitivity, or label fixes.
- Do not upload the full proof stack by default.
- Lead with the role-relevant proof: written answer, operating artifact, evidence log, RPE/profitability proof, sales/delivery automation proof, AI disclosure, and candidate profile.
- Treat HeavisideOS and Hermes as appendix-only proof if the portal supports a clean appendix; otherwise keep them for walkthrough or interview follow-up.
- Keep x-accel, decision council, and SloaneVault raw proof as backup unless requested.

The main remaining risk is not substance. It is attachment bloat: too many internal tool names can make strong evidence feel like an internal catalog.

Proof Artifact Upgrade

The new attachment set improves proof quality because it separates four different proof jobs:

Proof job	Artifact	Scoring effect
GM economic proof	`heaviside-rpe-profitability-redacted-summary.md`	Moves the role-level leverage claim into Tier 4 before/after data.
Measured AI workflow proof	`heavisideos-slack-eval-redacted-summary.md`	Shows benchmark, remediation, gate improvement, and retrieval lift with real eval counts.
Built-system proof	`heaviside-sales-delivery-automation-redacted-summary.md`	Shows real sales/delivery systems and source-record basis; time-savings claims remain estimate-labeled.
Cron-fleet operating proof	`hermes-cron-fleet-redacted-summary.md`	Shows Chris operates a monitored `88`-job Hermes scheduler surface with routing, status, pause/error visibility, and fleet health checks; use as bounded Tier 3 proof, not as proof every cron is valuable.

This materially reduces the risk that the packet reads like a polished strategy doc. It now has role-level economic proof, AI workflow proof, software-building proof, and bounded cron-fleet operating proof.

Baseline Comparison

Claude is competent on triage, risk tiering, and cron audit. The v2 packet should beat it because it adds:

- real revenue/FTE and margin proof,
- real sales and delivery automation examples,
- redacted proof artifacts backing the strongest claims,
- a more inspectable artifact,
- a source-labeled evidence log,
- a real self-assessment,
- explicit founder-leverage and revenue/FTE framing tied to the GM posting.
- public-source specificity from Single Grain / Single Brain / ClickFlow / Karrot instead of generic agency-ops recommendations.
- explicit wiring complement: Chris's INTP / `6w5`, strong Type `1`, Kolbe `9-4-3-2` profile as the follow-through/source-checking counterweight to Eric's ENTJ / `8w7`, Quick Start `9`, Follow Thru `2` pattern.
- bounded cron-fleet operating evidence from Hermes/open-agents rather than only a proposed cron-audit framework.
- sharper operating-experience wording around replacing manual agency drag with source-data -> generated deliverable -> operator review -> customer send/publish -> outcome logging.
- explicit talent-leadership and founder-leverage wording: agency sustained and grew while Chris held a full-time PASCO role, and current businesses take less than `5` hours/week to manage, backed by hiring competent operators, ownership, cadence, and selective intervention.

Claude Baseline Gap

| Challenge area | Claude baseline | Current packet |

|---|---|---|

| First 48 hours | Competent triage and stakeholder plan. | Similar triage quality, but stronger CEO-load design, explicit queue states, source labels, and no-unowned-item standard. |

| Risk tiering | Good green/yellow/red categories and examples. | Stronger because value/risk are separated, red flags override value, and failures update memory/rules. |

| Cron audit | Good sample judgments. | Stronger because it audits accepted work, review debt, source quality, owner clarity, and ties to Chris's own Hermes scheduler-fleet experience. |

| Operating experience | Claude explicitly lacks real examples. | Strong human advantage: Heaviside, HG Market Report, HG SEO Commander, adoption example, RPE/margin history, PASCO parallel operating proof. |

| MMPI/self-assessment | Claude cannot do this. | Strong human advantage: actual Kolbe/Enneagram/MBTI profile tied to Eric's operating gaps. |

| Evidence/proof | Claude has no proof artifacts. | Strong advantage: evidence log, RPE/margin summary, system summaries, eval proof, org chart, AI disclosure. |

Current read: this should beat Claude decisively on Part 4 and proof quality, and at least match or beat it on Parts 1-3.

Remaining Non-Scoring Blockers

1. Applicant should confirm the LinkedIn URL in `tasks/final-upload-v2/candidate-profile.md`.
2. Applicant should approve the self-assessment wording.
3. Applicant should approve the lean upload set versus appendix set in `tasks/final-upload-v2/attachment-map.md`.
4. Applicant should approve the existing-businesses framing: delegated operating capacity and explicit boundaries, not divided attention.
5. Re-run live posting check immediately before actual upload if time passes.

Verdict

The v2 packet is now materially stronger than Claude's published baseline and aligned with the public scoring guide. It is not merely a better essay; it has an inspectable operating artifact, a proof-tiered evidence log, and real operating metrics.