

Ops COO 009 - Written Answer

Thesis

The GM job is not to make Single Grain produce more AI output. It is to turn AI output into leverage: accepted work, fewer CEO approvals, better retention, stronger margins, and higher revenue per employee.

Single Grain's public positioning is already clear: do not advise on AI, run it inside the operating stack; tie work to pipeline and revenue, not vanity output [Observed from public websites]. I would apply that same standard internally.

I have been building that pattern inside my own companies. I started Heaviside in 2011 [Observed applicant claim] and have spent the last 2 years [Observed applicant claim] re-engineering Heaviside, Paving Marketers, and Garage Door Marketers around an intelligence layer, software-driven workflows, and CLI-driven operations. On a QuickBooks P&L basis, with FTE estimated from labor spend, Heaviside moved from \$55.6K revenue/FTE and 2.6% net margin in 2023 to \$77.2K and 14.2% in 2024, \$90.0K and 12.9% in 2025, and \$120.1K and 17.1% in 2026 Q1 [Observed internal P&L / Estimated FTE]. That is not the Single Grain target yet. It is proof that I understand the metric and can move it while business health improves.

My rule at Single Grain: an agent output does not count until it becomes an accepted action, explicit rejection, owner assignment, or rule/memory update.

First 48 Hours

I would not start with stakeholder interviews. I would start with a command queue.

- Hour 0-2 [Assumed]: snapshot Mission Control, Slack, CRM/HubSpot/Salesforce, Gong, GA4/GSC, ClickFlow, Oracle output, and expired outreach into one queue. Every one of the 23 pending items [Observed from brief] becomes dead, needs_source_check, green_auto, yellow_owner, yellow_council, or red_eric. - Hour 2-5 [Assumed]: handle ClickFlow churn first. The 2 lost customers this month [Observed from brief] are live retention damage in the TES/SaaS side of the business. I pull product usage, support history, contract value, churn signal, save path, and source links. Eric only sees it if the save requires pricing, a strategic relationship call, or customer-trust judgment. - Hour 5-7 [Assumed]: handle the 3 stale prospects [Observed from brief]. I check CRM stage, last reply, source, account fit, relationship risk, and owner, then assign a re-entry note with writeback or mark the item dead. No re-queueing. - Hour 7-10 [Assumed]: triage the backlog. Age matters, but content decides priority. By end of day 1 [Assumed], every item has an owner, rejection reason, or slot in Eric's decision block. - Day 2 [Assumed]: source-check Oracle's 15 quick-win keywords [Observed from brief], meet Shaun to understand agent failure modes, and end with no unowned yellow/red items.

Eric's 30 minutes [Observed from brief] are not for discovery. I need priority confirmation, never-auto-approve rules, what I can decide for 7 days [Assumed], decisions on only 3-5 red items [Observed from brief], and tomorrow's CEO leverage metric.

The brief says Eric starts 10 things and finishes 3 [Observed from brief]. I would not try to change that wiring. I would use it: Eric gets to create, sell, judge, and set direction; the operating system turns open loops into owners, dates, decisions, and outcomes.

Green / Yellow / Red Gate

I separate value from risk. A high value score never overrides a red flag.

Green auto-executes when value is ≥ 7 [Assumed], source is verified, the action is reversible, the pattern is known, writeback is available, and there are no red flags. Examples: Alfred internal briefings, approved-template Flash repurposing, and existing-page Oracle fixes with verified GSC/GA4 data.

Yellow goes to me, an owner, or a small council when judgment matters but CEO context does not: churn investigation, new-page SEO recommendations, non-ICP prospects, creative refreshes, channel pivots, or paid media restructuring.

Red goes to Eric: pricing, refunds, contracts, strategic accounts, public claims, legal/privacy/security, senior people decisions, new core crons, killing core crons, and CEO-level priority tradeoffs.

The value score uses revenue proximity, time sensitivity, source confidence, owner clarity, effort size, and learning value on a 0/1/2 scale [Assumed]. Risk flags are irreversibility, strategic relationship, brand/public exposure, legal/privacy/security, financial exposure, novelty, and low source confidence. When green fails, reverse it if possible, notify the owner, move that item type to yellow, log the failure, and change the rule.

Cron Audit

At day 30 [Observed from brief], I would audit the 48 crons [Observed from brief] by accepted work, not activity: action rate, revenue proximity, review burden, source quality, owner clarity, false-positive cost, false-negative cost, and memory feedback.

| Cron | Call | Reason | |---|---|---| | Competitor Pricing Monitor | Kill/pause | 3% action [Observed from brief]; daily pricing noise creates executive distraction. Replace with monthly or exception-based review [Assumed]. | | Candidate Outreach Follow-up | Kill/redesign | 5% response [Observed from brief]; low return plus brand risk. Replace with weekly top-5 human review [Assumed]. | | Churn Signal Detector | Improve urgently | 20% investigated [Observed from brief]; false negatives are expensive. Same-day yellow escalation with owner, source links, customer state, and save play. | | GSC Quick Win Scan | Improve | 12% actioned [Observed from brief]; surface only source-verified, existing-page, commercial-intent opportunities with clear owner and low effort. | | LinkedIn Prospect Sourcer | Improve | 8% meeting conversion [Observed from brief]; tighten ICP, trigger, and relationship filters, and batch daily. | | Deal Revival Scan | Keep | 60% followed up [Observed from brief]; strong revenue-proximate signal, but require outcome logging. | | Weekly Content Calendar | Keep/tighten | 50% used [Observed from brief]; feed publish/winner data back into memory. | | Content Repurpose - YouTube | Keep/tighten | 35% published [Observed from brief]; rejected hooks become training data. |

Operating Edge

I have built a lot of internal tools; two clean examples are HG Market Report and HG SEO Commander. HG Market Report replaced manual data retrieval and slide deck assembly with a live customer report, sales presentation, and follow-up sequence. It saves roughly 1-2 hours per sales report [Estimated applicant claim] and gives prospects a self-serve path when they want the analysis but do not want to talk to sales.

HG SEO Commander does the same thing on delivery: SEO analysis, reporting, GBP posts, blog posts, link-building deliverables, and reporting. It saves at least 3-4 hours per client per month [Estimated applicant claim]. The adoption lesson was simple: people use tools that remove bad work. When a bug created bad GBP posts, someone started deleting them manually. I fixed it through the CLI and showed them how to use the tooling instead. They were excited because the new process made them more capable.

The operating rhythm I installed was source data -> generated analysis/deliverable -> operator review -> customer send/publish -> outcome log. It stuck because the system removed manual deck/report work instead of adding ceremony.

For Oracle's 12 quick-win keywords from last night [Observed from brief], I would check source data, existing URL fit, search intent, expected lift, effort, owner, page sensitivity, and prior recommendation outcomes. Source-verified existing-page changes go green. New pages, mixed intent, or strategic positioning go yellow. Public claims, strategic pages, or customer-sensitive pages go red.

The GM in 2028 [Observed from brief] owns the boundary between automation and judgment: what executes alone, what routes to an owner, what goes to expert review, what stays with the CEO, and how failures become better rules. Humans keep trust, pricing, strategic accounts, people decisions, public claims, and founder-level judgment. Systems handle coordination, source checks, repetitive analysis, routing, reporting, and memory.

My self-assessment: I am a founder first. My current shorthand is INTP / 6w5 with a strong Type 1 overlay, Kolbe 9-4-3-2 [Observed self-assessment]. The practical read is high source-checking, preparedness, standards, error correction, and lower appetite for random churn. That is a useful counterweight to an ENTJ / 8w7, Quick Start 9, Follow Thru 2 CEO [Observed from brief]. I am not another idea generator. I am the person who turns 10 starts into 10 owned outcomes.

The closest operating analogy is forward-deployed engineer inside my own companies: I sit inside the business problem, build the tool, deploy it, and own whether it works. For competing workstreams, my mechanism is GTD/ToDoist discipline plus SloaneVault as a CLI-drivable extended brain. The risk is overbuilding or pushing capability people do not need, want, or use. I compensate by fighting capability sprawl and forcing tools back to accepted work, owner action, adoption, and measured outcomes.

This breaks if source data is wrong, yellow items have no owner, Eric keeps reviewing non-red decisions, memory does not get updated, councils become delay theater, or model/tool spend is unmanaged. The artifact shows the queue states, gate, cron audit, cadence, memory schema, human boundary, and failure modes.