

Ops COO 009 - Operational Takeover Plan

LinkedIn: <https://www.linkedin.com/in/christophersloane/>

Numbers note: numbers from the challenge brief are observed from the brief. My company numbers are from internal records. Time-savings claims are estimates. Working time blocks are assumptions.

AI output is cheap. The job is turning it into accepted work: action taken, owner assigned, rejected with a reason, or fed back into memory so the next pass is better.

That is how I would run this. Not more dashboards. Not more meetings. A simple queue, a risk gate, a daily decision lane, and fewer approvals for Eric.

I have been doing this in my own companies. Internal records: I started Heaviside in 2011 and still run it, along with Paving Marketers and Garage Door Marketers. Internal P&L/FTE model: revenue per FTE moved from \$55.6K with 2.6% net margin in 2023 to \$120.1K with 17.1% net margin in 2026 Q1. The company got simpler after an overstaffed period, but this was not artificial shrinkage. The point was more useful work per person.

1. First 48 hours

First rule: stop revenue leaks before cleaning up the queue.

Assumed hours 0-2: pull Mission Control, Slack, CRM, Gong, GA4/GSC, ClickFlow, Oracle output, and expired outreach into one command queue. The observed 23 pending items get one state: dead, source-check, green, yellow, or red.

Assumed hours 2-5: handle ClickFlow churn first. The brief says two customers were lost this month. I would pull usage, support history, plan, contract value, last touch, churn reason, and possible save action. Eric only sees it if the save requires pricing, a relationship call, or customer-trust judgment.

Assumed hours 5-7: handle the three stale prospects. Check stage, last reply, deal value, owner, relationship risk, and next best action. If recoverable, sales gets a re-entry note and a deadline. If not, close it. No re-queueing.

Assumed hours 7-10: clear the remaining backlog by type. Green items execute. Yellow items get an owner and due time. Red items go into Eric's decision block. By end of day one, every item has an owner, a rejection reason, or a CEO decision slot.

Day two: source-check Oracle's observed 15 quick-win keywords, meet Shaun to understand agent failure modes, and set the daily operating lane. Eric's 30 minutes are for five decisions: top priority, never-auto-approve rules, my seven-day decision rights, the highest-risk red items, and tomorrow's CEO-time metric.

The brief says Eric starts 10 things and finishes 3. I would not try to change his wiring. Eric should create, sell, judge, and set direction. The operating system should turn open loops into owners, dates, decisions, and outcomes.

2. Risk gate

The gate is simple: value gets scored, risk gets veto power.

Green auto-executes when the source is verified, the action is reversible, the pattern is known, the owner or writeback path is clear, and there are no red flags. Examples: internal Alfred briefings, approved-template Flash repurposing, and existing-page Oracle updates backed by GSC/GA4.

Yellow goes to me or an accountable owner. This is for judgment that does not need Eric: churn investigation, new-page SEO recommendations, non-ICP prospects, creative refreshes, channel pivots, and paid-media restructuring.

Red goes to Eric. Pricing, refunds, contracts, strategic accounts, public claims, legal/privacy/security, senior people decisions, new core crons, killed core crons, and CEO-level priority tradeoffs.

The assumed target is 3-5 Eric items per day. If green fails, reverse it if possible, notify the owner, move that item type to yellow, log the miss, and change the rule. One miss is learning. A repeat miss means the gate is wrong.

3. Cron audit

At day 30, I would audit the observed 48 crons by accepted work, not output volume. The test is: did someone act, did the action matter, and could we get the same result with less noise?

Kill or pause two. Competitor Pricing Monitor has an observed 3% action rate. Daily pricing noise is not worth executive attention; make it monthly or exception-based. Candidate Outreach Follow-up has an observed 5% response rate and creates brand risk; replace it with a weekly human review of the top five stale candidates.

Improve three. GSC Quick Win Scan has an observed 12% action rate, so tighten it to source-verified, existing-page, commercial-intent changes with a clear owner. LinkedIn Prospect Sourcer has an observed 8% meeting conversion, so reduce volume and tighten ICP, trigger, and relationship filters. Churn Signal Detector has an observed 20% investigation rate; make it same-day yellow escalation with usage data, owner, save play, and outcome logging.

Keep and tighten three. Deal Revival Scan has an observed 60% follow-up rate, so keep it and require outcome logging. Weekly Content Calendar has an observed 50% use rate, so keep it but feed publish data back into memory. Content

Repurpose has an observed 35% publish rate, so keep it with stronger rejection feedback.

This is familiar ground for me. I operate Hermes/open-agents with 80+ scheduled jobs across research, content, inbox, infrastructure, and cold email. The lesson is not that every job should live. Jobs need owners, pause states, error visibility, and accepted-work feedback.

4. Operating edge

Two examples from my own agency: HG Market Report and HG SEO Commander.

HG Market Report moved sales audits out of manual data pulls and slide decks. It captures prospect context, queues local SEO/SERP/geogrid checks, builds the analysis, creates a public report URL, and sends the report-ready email. Estimated savings: 1-2 hours per standard sales audit.

HG SEO Commander is the delivery version. It pulls GA4, Search Console, Google Business Profile, SERP/geogrid tracking, citations, articles, GBP posts, approvals, scheduled syncs, queues, and admin tools into one operating surface. Estimated savings: 3-4 hours per active SEO client per month. Strategy and quality review stay human.

The before state was normal agency drag: manual data retrieval, report assembly, ad hoc review, and work living in people's heads. The after state is source data, generated analysis or deliverable, operator review, customer send or publish, and outcome log. When a GBP-post bug created bad posts, someone started deleting them manually. I fixed it through the CLI and showed the operator the cleaner path. The process stuck because the tool removed bad work.

To review Oracle's 12 quick-win keywords, I would check source data, existing URL fit, intent, expected lift, effort, owner, page sensitivity, and prior recommendation history. Existing-page, source-verified changes can go green. New pages, mixed intent, or positioning calls go yellow. Public claims or strategic pages go red.

A 2028 GM in this company owns the boundary between automation and judgment. Systems handle retrieval, scoring, routing, reporting, drafts, and memory. Humans keep pricing, trust, strategic accounts, people decisions, public claims, and final judgment.

AI disclosure: I used Claude/Codex, ChatGPT-style review, SloaneVault, x-accel-style scoring, and council-style critique. AI helped me read the brief, compare against the rubric and Claude baseline, pressure-test the plan, and tighten the writing. I made the operating calls, checked the source documents, and removed generic AI language. Some proof is private. Public summary: <https://csloane.com/numbers>. Raw records can be walked through if requested.

Self-assessment: my shorthand is INTP, likely 6w5 with a strong Type 1 overlay, Kolbe 9-4-3-2. In practice that means source-checking, preparation, standards, and error correction. That pairs well with an ENTJ/8w7, Quick Start 9, Follow Thru 2 CEO. I am not another idea generator. I am the person who turns 10 starts into 10 owned outcomes.

This breaks if source data is wrong, yellow items have no owner, or Eric keeps reviewing non-red decisions. It also breaks if memory does not get updated or agents are measured by output instead of accepted work. The fix is not more process. The fix is closing the loop every day.

Evidence log, short version: P&L/FTE claims are private before/after records. HG Market Report and HG SEO Commander are inspectable software/workflow artifacts. Hermes is backed by logs and job manifests. Personality is self-assessment. The public version of the proof lives on /numbers; confidential records stay out of the initial application.