

Ops COO 009 - Written Answer

Thesis

The GM job is not to make Single Grain produce more AI output. It is to turn AI output into leverage: accepted work, fewer CEO approvals, better retention, stronger margins, and higher revenue per employee.

Single Grain's public positioning is already clear: do not advise on AI, run it inside the operating stack; tie work to pipeline and revenue, not vanity output [Observed from public websites]. I would apply that same standard internally.

The strategic lever is the goal of 50% TES revenue by 2028 [Observed from brief]: hold core services operationally tight while ClickFlow, Karrot, and agent-enabled workflows turn outputs into productized revenue, so revenue/FTE rises through mix, productivity, and accepted work, not artificial shrinkage.

I have been building that pattern inside my own companies. I started Heaviside in 2011 [Observed applicant claim] and have spent the last 2 years [Observed applicant claim] re-engineering Heaviside, Paving Marketers, and Garage Door Marketers around an intelligence layer, software-driven workflows, and CLI-driven operations. On a QuickBooks P&L basis, with FTE estimated from labor spend, Heaviside moved from \$55.6K revenue/FTE and 2.6% net margin in 2023 to \$77.2K and 14.2% in 2024, \$90.0K and 12.9% in 2025, and \$120.1K and 17.1% in 2026 Q1 [Observed internal P&L / Estimated FTE]. That is not the Single Grain target yet. It is proof that I understand the metric and can move it while business health improves.

The people side is just as important. I have sustained and grown the agency while holding a full-time role at PASCO, where I own SEO/digital marketing and now ML/AI work [Observed applicant claim]. My existing businesses take less than 5 hours/week to manage [Observed applicant claim] because I hire competent operators, give them ownership, and do not micromanage. That is the leadership model I would bring here: hire people who can own outcomes, give them a scoreboard and cadence, intervene quickly when ownership fails, and keep the CEO out of routine resolution.

My rule at Single Grain: an agent output does not count until it becomes an accepted action, explicit rejection, owner assignment, or rule/memory update.

First 48 Hours

I would not start with stakeholder interviews. I would start with a command queue.

- Hour 0-2 [Assumed]: snapshot Mission Control, Slack, CRM/HubSpot/Salesforce, Gong, GA4/GSC, ClickFlow, Oracle output, and expired outreach into one queue. Every one of the 23 pending items [Observed from brief] becomes dead, needs_source_check, green_auto, yellow_owner, yellow_council, or red_eric. - Hour 2-5 [Assumed]: handle ClickFlow churn first. The 2 lost customers this month [Observed from brief] are live retention damage in the TES/SaaS side of the business. I pull product usage, support history, contract value, churn signal, save path, and source links. Eric only sees it if the save requires pricing, a strategic relationship call, or customer-trust judgment. - Hour 5-7 [Assumed]: handle the 3 stale prospects [Observed from brief]. I check CRM stage, last reply, source, account fit, relationship risk, and owner, then assign a re-entry note with writeback or mark the item dead. No re-queueing. - Hour 7-10 [Assumed]: triage the backlog. Age matters, but content decides priority. By end of day 1 [Assumed], every item has an owner, rejection reason, or slot in Eric's decision block. - Day 2 [Assumed]: source-check Oracle's 15 quick-win keywords [Observed from brief], meet Shaun to understand agent failure modes, and end with no unowned yellow/red items.

Eric's 30 minutes [Observed from brief] are not for discovery. I need priority confirmation, never-auto-approve rules, what I can decide for 7 days [Assumed], decisions on only 3-5 red items [Observed from brief], and tomorrow's CEO leverage metric.

The brief says Eric starts 10 things and finishes 3 [Observed from brief]. I would not try to change that wiring. I would use it: Eric gets to create, sell, judge, and set direction; the operating system turns open loops into owners, dates, decisions, and outcomes.

Green / Yellow / Red Gate

I separate value from risk. A high value score never overrides a red flag.

Green auto-executes when value is ≥ 7 [Assumed], source is verified, the action is reversible, the pattern is known, writeback is available, and there are no red flags. Examples: Alfred internal briefings, approved-template Flash repurposing, and existing-page Oracle fixes with verified GSC/GA4 data.

Yellow goes to me, an owner, or a small council when judgment matters but CEO context does not: churn investigation, new-page SEO recommendations, non-ICP prospects, creative refreshes, channel pivots, or paid media restructuring.

Red goes to Eric: pricing, refunds, contracts, strategic accounts, public claims, legal/privacy/security, senior people decisions, new core crons, killing core crons, and CEO-level priority tradeoffs.

The value score uses revenue proximity, time sensitivity, source confidence, owner clarity, effort size, and learning value on a 0/1/2 scale [Assumed]. Risk flags are irreversibility, strategic relationship, brand/public exposure, legal/privacy/security, financial exposure, novelty, and low source confidence. When green fails, reverse it if possible, notify the owner, move that item type to yellow, log the failure, and change the rule.

Cron Audit

At day 30 [Observed from brief], I would audit the 48 crons [Observed from brief] by accepted work, not activity: action rate, revenue proximity, review burden, source quality, owner clarity, false-positive cost, false-negative cost, and memory feedback.

| Cron | Call | Reason | |---|---|---| | Competitor Pricing Monitor | Kill/pause | 3% action [Observed from brief]; daily pricing noise creates executive distraction. Replace with monthly or exception-based review [Assumed]. | | Candidate Outreach Follow-up | Kill/redesign | 5% response [Observed from brief]; low return plus brand risk. Replace with weekly top-5 human review [Assumed]. | | Churn Signal Detector | Improve urgently | 20% investigated [Observed from brief]; false negatives are expensive. Same-day yellow escalation with owner, source links, customer state, and save play. | | GSC Quick Win Scan | Improve | 12% actioned

[Observed from brief]; surface only source-verified, existing-page, commercial-intent opportunities with clear owner and low effort. | | LinkedIn Prospect Sourcer | Improve | 8% meeting conversion [Observed from brief]; tighten ICP, trigger, and relationship filters, and batch daily. | | Deal Revival Scan | Keep | 60% followed up [Observed from brief]; strong revenue-proximate signal, but require outcome logging. | | Weekly Content Calendar | Keep/tighten | 50% used [Observed from brief]; feed publish/winner data back into memory. | | Content Repurpose - YouTube | Keep/tighten | 35% published [Observed from brief]; rejected hooks become training data. |

This is not theoretical for me. I operate a Hermes/open-agents scheduler fleet; on 2026-05-14, repo verification checked 88 jobs, 28 skills, and 3 managed contexts with no job-manifest issues, and live fleet status reported healthy across 3/3 nodes [Observed internal scheduler state]. The lesson is not that every cron should live; it is that crons need owners, pause/kill states, error visibility, and accepted-work feedback.

Operating Edge

I have built a lot of internal tools; two clean examples are HG Market Report and HG SEO Commander. HG Market Report turns a prospect audit into a software workflow: capture the business and Google Business Profile context, generate keywords and service area, queue SEO audit / SERP / geogrid reports, derive lost-revenue / quick-win / competitor views, create a public presentation URL, and send the report-ready email. That replaces the manual research and deck loop and conservatively saves 1-2 hours per standard sales audit [Estimated applicant claim].

HG SEO Commander is the delivery version: GA4, Search Console, Google Business Profile, SERP/geogrid tracking, citations, AI articles, GBP posts, bulk-posting approvals, scheduled syncs, queues, and CLI/admin operations. It conservatively removes 3-4 hours per active SEO client per month from routine data pulls, status checks, report assembly, GBP/content prep, and queue recovery [Estimated applicant claim]. Strategy and quality review stay human. The adoption lesson was simple: people use tools that remove bad work. When a bug created bad GBP posts, someone started deleting them manually. I fixed it through the CLI and showed them how to use the tooling instead. They were excited because the new process made them more capable.

The broken rhythm before was familiar agency drag: manual data pulls, slide assembly, ad hoc review, and operators carrying too much process in their heads. I replaced it with source data -> generated analysis/deliverable -> operator review -> customer send/publish -> outcome log. Enforcement was not another meeting; it was making the tool the easiest path to good work. It stuck because the system removed bad work instead of adding ceremony.

For Oracle's 12 quick-win keywords from last night [Observed from brief], I would check source data, existing URL fit, search intent, expected lift, effort, owner, page sensitivity, and prior recommendation outcomes. Source-verified existing-page changes go green. New pages, mixed intent, or strategic positioning go yellow. Public claims, strategic pages, or customer-sensitive pages go red.

The GM in 2028 [Observed from brief] owns the boundary between automation and judgment: what executes alone, what routes to an owner, what goes to expert review, what stays with the CEO, and how failures become better rules. Humans keep trust, pricing, strategic accounts, people decisions, public claims, and founder-level judgment. Systems handle coordination, source checks, repetitive analysis, routing, reporting, and memory.

My self-assessment: I am a founder first. My current shorthand is INTP / 6w5 with a strong Type 1 overlay, Kolbe 9-4-3-2 [Observed self-assessment]. The practical read is high source-checking, preparedness, standards, error correction, and lower appetite for random churn. That is a useful counterweight to an ENTJ / 8w7, Quick Start 9, Follow Thru 2 CEO [Observed from brief]. I am not another idea generator. I am the person who turns 10 starts into 10 owned outcomes.

The closest operating analogy is forward-deployed engineer inside my own companies: I sit inside the business problem, build the tool, deploy it, and own whether it works. For competing workstreams, my mechanism is GTD/ToDoist discipline plus SloaneVault as a CLI-drivable extended brain. The risk is overbuilding or pushing capability people do not need, want, or use. I compensate by fighting capability sprawl and forcing tools back to accepted work, operator adoption, and measured outcomes.

This breaks if source data is wrong, yellow items have no owner, Eric keeps reviewing non-red decisions, memory does not get updated, councils become delay theater, or model/tool spend is unmanaged. The artifact shows the queue states, gate, cron audit, cadence, memory schema, human boundary, and failure modes.

Appendix / Supporting Artifacts

Not part of the written-answer page count.

The official written answer is the first 2 pages of this packet.

The following pages are image-rendered supporting artifacts: operating artifact, evidence log, source labels, AI disclosure, candidate profile, proof summaries, application support, and internal QA.

The challenge brief says diagrams and decision trees do not count toward the limit. The public scoring guide says artifact links, code, sheets, and short demos do not count toward the written page limit unless the brief says otherwise.

If the application portal counts every PDF page mechanically, submit the 2-page official PDF and attach or link supporting artifacts separately.

ai-usage-disclosure

Ops COO 009 - AI Usage Disclosure v2

Tools Used

- Codex / ChatGPT-style assistant for research organization, critique, drafting, compression, and QA.
- Local repo inspection for the challenge brief, `SCORING.md`, `ops-coo-009/scoring\_rubric.md`, and `ops-coo-009/claude\_baseline.md`.
- SloaneVault/qmd search for internal source records, project notes, Eric/Single Grain research, x-accel references, HeavisideOS references, and proof candidates.
- Web research for official Single Grain GM role facts and source links.
- x-accel authenticated X capture stored in SloaneVault for Eric's public commentary, refreshed on 2026-05-13.
- Hermes/open-agents scheduler-state review for bounded cron-fleet proof.
- Decision-council review and peer-review synthesis to stress-test the packet before final packaging.
- Local PDF export / QA scripts for page-fit, number-label review, stale-claim scan, and sensitive-info scan.

What AI Helped With

- Comparing the draft against Claude's published baseline.
- Organizing the answer around the GM job mandate, not just the challenge mechanics.
- Drafting and compressing the written answer, operating artifact, evidence log, and disclosure.
- Stress-checking source labels, proof tiers, failure modes, human boundaries, and score-losing claims.
- Separating HeavisideOS, x-accel, SloaneVault, and decision councils accurately.
- Identifying the need to add bounded Hermes scheduler-fleet evidence without overclaiming cron quality.

What I Personally Decided

- The core thesis: AI output only matters when it becomes accepted work, owner action, memory, margin, retention, and revenue/FTE leverage.
- The first 48-hour priority order: ClickFlow churn first, stale prospects second, queue routing third, Oracle keyword source checks fourth.
- The green/yellow/red routing model: green only for reversible, source-verified, known-pattern work; yellow for owner/council judgment; red for Eric-level trust, pricing, strategy, public, legal, people, and key relationship decisions.
- The cron audit posture: kill or pause low-action/high-noise workflows instead of keeping every automation alive.
- The evidence posture: lead with Heaviside RPE/profitability proof, support with HG Market Report, HG SEO Commander, SloaneVault, x-accel, decision councils, and carefully bounded HeavisideOS proof.
- The cron proof posture: use Hermes/open-agents as evidence of operating a monitored scheduler fleet, not as proof that every scheduled job deserves to exist.
- The maturity boundary: x-accel is separate from HeavisideOS; HeavisideOS is early practical Slack/client/performance summarization plus company-brain direction, not a fully adopted company operating system.
- The redaction rule: raw internal reports, P&L exports, CRM data, Slack messages, and private vault notes should not be attached externally without redaction.

What I Checked

- `challenges/ops-coo-009/brief.md`
- `challenges/ops-coo-009/scoring\_rubric.md`
- `SCORING.md`
- `challenges/ops-coo-009/claude\_baseline.md`
- GM role posting facts for compensation and remoteness.
- Heaviside RPE/profitability calculations from QuickBooks P&L snapshots and labor-to-FTE method.
- Applicant interview notes for HG Market Report, HG SEO Commander, operating style, self-assessment, and approved wording.
- SloaneVault project notes and lessons for corrections and maturity boundaries.
- Number labels in the v2 written answer and artifact.
- Sensitive-info risk in internal proof artifacts.
- Hermes scheduler manifests and live-state counts used in the redacted cron-fleet summary.

Known Limits

- GM role facts should be rechecked immediately before upload because public job pages can change.
- Eric/X commentary is based on captured corpora refreshed through 2026-05-13 17:52 UTC, not a guarantee of later posts.
- Heaviside RPE/FTE history uses observed P&L revenue and net income, with FTE estimated from labor spend using documented cost/FTE benchmarks.
- HG Market Report and HG SEO Commander time savings are applicant estimates unless later backed by time logs.
- Hermes/open-agents scheduler proof is internal source-record evidence and does not independently verify business value for every scheduled job.
- Internal proof summaries should be attached only in redacted form.

application-answers

Application Questions - Working Answers

Status: draft for applicant approval.

These answers should stay consistent with the challenge packet, cover letter, and proof artifacts.

1. Have you worked at a digital marketing agency before?

Short answer:

> Yes. I founded Heaviside in 2011 and have operated it since then across SEO, web design, digital marketing, and related services, including the Paving Marketers and Garage Door Marketers vertical niches.

Expanded answer:

> Yes. I founded Heaviside in 2011 and have operated it since then. The agency started with SEO, web design, and digital marketing, then expanded into related services including cold email, direct mail, social media, and verticalized offers through Paving Marketers and Garage Door Marketers. I am still the founder/operator, but the agency no longer depends on daily founder involvement. My existing businesses take less than '5' hours/week to manage because I have hired competent operators, given them ownership, and built more of the work into systems.

2. Can you walk me through how you currently use AI in your day-to-day work?

Short answer:

> I use AI as an operating layer, not just a writing assistant: source gathering, analysis, coding, workflow automation, Slack/client-performance summaries, research, content drafting, scoring gates, and agent/crons that produce work only when there is a route to owner action or outcome logging.

Expanded answer:

> My daily AI use is very practical. I use Codex/Claude Code-style tools for development, debugging, repo analysis, and CLI-driven workflow work. I use SloaneVault as a durable memory layer so decisions, sources, project state, and lessons do not disappear between sessions. I use HeavisideOS in its current practical form for SlackBot summaries and query flows around recent client communication and paid ads/SEO/business-performance context. I use x-accel-style scoring gates to decide whether content or agent output is ready to move forward. I also operate Hermes/open-agents, an agent and scheduled-job layer that handles research, content drafting, infrastructure checks, project status, inbox/briefing flows, fleet monitoring, and cold-email operations. The important rule is that AI output does not count until it becomes accepted work, an explicit rejection, owner action, a system-of-record update, or a rule/memory update.

3. Can you give an example of a process or system you built that permanently improved execution?

Short answer:

> HG SEO Commander is the cleanest delivery example. It moved recurring SEO work from manual data retrieval, reporting, content/GBP prep, and queue recovery into a software workflow with source data, scheduled syncs, queues, CLI/admin tools, operator review, and customer-ready deliverables.

Expanded answer:

> One example is HG SEO Commander. Before it, recurring SEO delivery depended on operators manually pulling data from multiple sources, assembling reports, checking rankings/local visibility, and producing deliverables like GBP posts and blog posts. The system now integrates GA4, Search Console, Google Business Profile, SERP/geogrid tracking, citations, AI articles, GBP posts, bulk posting approvals, scheduled syncs, queues, and CLI/admin operations. It conservatively removes '3-4' hours per active SEO client per month from routine data pulls, status checks, report assembly, GBP/content prep, and queue recovery. Strategy and quality review stay human, but the repetitive execution layer moved into software. The adoption lesson was important: people use tools that remove bad work. When a bad GBP-post bug created manual cleanup, I fixed it through the CLI and showed the operator how to use the tooling instead. They were excited because the new process made them more capable.

Alternate sales-side example:

> HG Market Report automated the front-end sales audit. Instead of manually collecting business/GBP context, pulling local search data, taking screenshots, building a slide deck, and sending a custom report, the system captures prospect context, generates keywords and service area, queues SEO audit / SERP / geogrid reports, derives lost-revenue / quick-win / competitor views, creates a public presentation URL, and sends the report-ready email. It saves an estimated '1-2' hours per standard sales audit and gives prospects a useful path even when they do not want to engage sales immediately.

Application Consistency Notes

- The "MMPI-style" item is the challenge's self-assessment request, not a clinical MMPI. Our packet uses LLM-guided self-assessment plus Kolbe / Enneagram / MBTI-style operating-profile work and compares the practical wiring to Eric's profile in the brief.
- Do not overclaim HeavisideOS. Current real use is SlackBot summaries/query flows around client communication and paid ads/SEO/business-performance context, plus the broader company-brain direction.
- Do not claim x-accel is part of HeavisideOS. It is separate.
- Do not present Hermes cold email as a clean current-performance proof unless live credentials/launch state are verified. Use it as built-system / closed-loop operating evidence.
- Keep the existing-businesses point framed as delegated operating capacity, not casual divided attention.

APPROVAL-CHECKLIST

Ops COO 009 - Approval Checklist

Reply with approvals or edits for the items below.

1. Official challenge answer

Approve `00\_OFFICIAL\_challenge\_answer\_2\_pages.pdf` as the official written answer.

Reason: this is the controlled 2-page response to the COO Operations Challenge.

2. Upload strategy

Approve this upload strategy:

- Use the 2-page PDF as the official written answer.
- Upload the operating artifact, evidence log, AI disclosure, candidate profile, and core proof summaries as supporting files if the portal allows it.
- Keep the full image-appendix PDF as a review/backup packet unless the portal clearly allows appendices that do not count toward the written limit.

3. Cover letter

Approve the cover letter in `cover-letter.md` for the application cover-letter field.

Key line: "I was not looking for a new role. This one found me."

4. Application answers

Approve the three application answers in `application-answers.md`:

- Agency experience
- Day-to-day AI use
- Process/system that permanently improved execution

5. Candidate profile

Approve the LinkedIn URL and self-assessment framing in `candidate-profile.md`.

Key framing: founder/operator, INTP / 6w5 with strong Type 1 overlay, Kolbe 9-4-3-2, complementary to Eric's briefed profile.

6. Evidence claims

Approve these major claims as accurate enough to submit:

- Heaviside started in 2011.
- Heaviside does SEO, web design, digital marketing, and related services.
- You are still founder/operator.
- Current Heaviside/PVM/GDM management load is less than 5 hours per week.
- Heaviside revenue/FTE improved materially while net margin improved, using internal P&L and estimated FTE methodology.
- HG Market Report saves an estimated 1-2 hours per standard sales audit.
- HG SEO Commander saves an estimated 3-4 hours per active SEO client per month.
- Hermes/open-agents evidence can be described as a real internal scheduler/agent fleet, not as a public SaaS claim.
- HeavisideOS should be described narrowly as current SlackBot summaries/query flows plus broader company-brain direction.
- x-accel is separate from HeavisideOS.

7. Public link

Approve using `https://csloane.com/numbers` as the public background link in the cover letter.

8. Redaction boundary

Approve using redacted proof summaries instead of raw financial/client/system exports.

Reason: the packet should support source checks and a live walkthrough without exposing unnecessary confidential detail in the initial application.

attachment-map

Ops COO 009 - Attachment Map

Use this as the plain-language reviewer guide if the portal allows multiple uploads, or as the cover page if the proof files are merged into one appendix.

Upload Hierarchy

The reviewer path should stay lean. Do not make the first reader sort through the full internal proof stack.

Tier	Files	Rule
Required answer	`written-answer.pdf`, `ai-usage-disclosure.md`, `candidate-profile.md`	Submit wherever the application explicitly asks for the challenge answer, AI usage disclosure, profile, or self-assessment.
Core proof	`operating-artifact.md`, `evidence-log.md`, `attachment-map.md`, `heaviside-rpe-profitability-redacted-summary.md`, `heaviside-sales-delivery-automation-redacted-summary.md`, `heaviside-org-chart-redacted-summary.md`	Include if the portal allows supporting files or if merging into one appendix. These are the strongest role-relevant proof files.
Application support	`cover-letter.md`, `application-answers.md`	Use for the application form and cover-letter field. Do not upload unless the portal asks for those fields/files.
Appendix only	`heavisideos-slack-eval-redacted-summary.md`, `hermes-cron-fleet-redacted-summary.md`	Include only if the upload surface supports a clean appendix. These strengthen AI/workflow and cron-fleet credibility, but should not distract from GM economic proof.
Backup / walkthrough	`job-post-line-scorecard.md`, `challenge-compliance-audit.md`, `hermes-closed-loop-agent-proof.md`, x-accel, decision council, SloaneVault memory, raw screenshots, raw exports	Hold for interview, public site, or follow-up unless requested. Do not attach raw confidential material.

If the portal is cramped, submit the required answer plus the core proof in this order: written answer, operating artifact, evidence log, RPE/profitability proof, sales/delivery automation proof, AI disclosure, candidate profile.

File	What it proves	Why it matters
`written-answer.pdf`	Main response to the COO Operations Challenge.	Required answer; fits the written cap.
`operating-artifact.md`	Queue states, risk gate, CEO operating contract, cron audit, dashboard, memory schema, cadence, failure modes.	Shows the operating system a reviewer can inspect.
`evidence-log.md`	Claim-to-proof mapping with source labels and proof tiers.	Keeps numbers and proof disciplined under `SCORING.md`.
`ai-usage-disclosure.md`	Tools used, what AI helped with, what Chris decided, and what was checked.	Required disclosure; separates AI help from human judgment.
`candidate-profile.md`	LinkedIn, MMPI-style operating profile, and founder/talent-leadership model.	Required profile/self-assessment material; also addresses leadership-bench fit.
`cover-letter.md`	Concise application cover letter.	Use in the cover-letter field if available.
`application-answers.md`	Draft answers for agency experience, day-to-day AI usage, and permanent process/system improvement.	Use for application form fields and interview consistency.
`heaviside-rpe-profitability-redacted-summary.md`	Heaviside revenue/FTE and net margin improved together from P&L/FTE model.	Primary GM economic proof.
`heaviside-sales-delivery-automation-redacted-summary.md`	HG Market Report and HG SEO Commander as real sales/delivery automation systems.	Built-system proof; time savings remain estimate-labeled.
`heaviside-org-chart-redacted-summary.md`	Public roster plus functional ownership map across technical operations, web, SEO, PPC, operations, client success, and reporting/admin.	Supports talent leadership, delegation, leadership bench, and founder-leverage claims. Not for FTE validation.
`heavisideos-slack-eval-redacted-summary.md`	Before/after Slack assistant eval and retrieval improvement.	Optional measured AI workflow proof; attach only if the packet can support one more proof file cleanly.
`hermes-cron-fleet-redacted-summary.md`	`88` repo-checked Hermes/open-agents jobs with routing, status, pause/error visibility, and fleet health checks.	Bounded cron-fleet operating proof; use as backup or walkthrough evidence.
`job-post-line-scorecard.md`	Line-by-line scoring against the GM job post's exclusion and strong-candidate language.	Internal QA; prevents generic scoring.
`challenge-compliance-audit.md`	Requirement-by-requirement audit against the challenge parts, win criteria, lose criteria, and public scoring standards.	Internal QA; confirms the packet follows the prompt exactly.
`hermes-closed-loop-agent-proof.md`	Source-backed summary of Hermes as an outputs-to-outcomes operating layer across research, content, infra, inbox, fleet, and cold-email jobs.	Backup proof / public-page source material.

candidate-profile

Candidate Profile Addendum

LinkedIn

<https://www.linkedin.com/in/christophersloane/>

MMPI-Style Self-Assessment Summary

Note: this is the "MMPI-style" item from the challenge brief, not a clinical MMPI-2/MMPI-3 instrument. For this application, I used LLM-guided self-assessment plus Kolbe / Enneagram / MBTI-style operating-profile work, then compared the practical operating implications against Eric's profile in the brief.

I am a founder first. My current shorthand is **INTP / `6w5`, Kolbe `9-4-3-2`**. The Enneagram result is not a clean landslide: Type `1` and Type `6` tied at `15`, with Type `3` at `12` and Type `5` at `11`. I call it `6w5` because the Type `6` pattern has a coherent Type `5` wing, while Type `1`'s usual wings, Type `9` and Type `2`, were basically absent. Practical read: trust/security/preparedness core, with a strong standards/responsibility/error-correction overlay.

Kolbe: Fact Finder `9`, Follow Thru `4`, Quick Start `3`, Implementor `2`. I naturally source-check, define terms, compare options, and preserve enough structure for the work to keep moving. I am not another high-Quick-Start idea generator.

That matters for this role because the brief describes Eric as ENTJ / `8w7`, Kolbe Quick Start `9`, Follow Thru `2`, and self-rated weak on follow-through, organization, and workflow consistency. I would not try to make him less like that. I would use his strengths and build the follow-through layer around him: clear owners, dates, decision logs, memory, and accepted-work tracking.

My strongest traits for this role are high agency, attention to detail, systems thinking, and tolerance for ambiguity. I am comfortable sitting inside a messy business problem, building the tool or operating layer needed to solve it, deploying it, and owning whether it works.

The leadership proof is that my agency does not require daily founder intervention. I have sustained and grown it while holding a full-time role at PASCO, where I own SEO/digital marketing and now ML/AI work. Today my existing businesses take less than `5` hours/week to manage because I hire competent people, give them ownership, and avoid micromanagement. It also matters for this role because I am not relying on founder heroics to keep my existing businesses alive; they are built around delegated ownership and operating leverage. My default is to hire people who can get the job done, give them a clear scorecard and cadence, and step in only where ownership, judgment, or performance actually requires it.

The closest operating analogy is forward-deployed engineer inside my own companies. I do not just specify systems; I build and run them against real sales, delivery, reporting, client communication, and fulfillment problems.

The main risk is overbuilding. I can create capability faster than a team can absorb it, or build things people do not need, want, or use. I compensate by fighting capability sprawl and forcing tools back to accepted work, owner action, adoption, and measured outcomes.

I work best with a founder who does not want to micromanage, hires competent people, expects them to get the job done, and steps in only where CEO judgment is truly needed. That is the operating fit I would aim for here: Eric keeps creating leverage; I make sure the company does not leak the follow-through.

challenge-compliance-audit

Ops COO 009 - Challenge Compliance Audit

Status: working QA.

Sources:

- `challenges/ops-coo-009/brief.md`
- `challenges/ops-coo-009/scoring\_rubric.md`
- `SCORING.md`
- `challenges/ops-coo-009/claude\_baseline.md`

Required Packet

Requirement	Current status	Score	Notes
Written answer	Complete	`9/10`	`written-answer.pdf` is `2` pages.
Operating artifact	Complete	`9.25/10`	Includes queue, risk gate, cron audit, dashboard, cadence, memory schema, failure modes, and now a more literal `48`-hour timeline.
Evidence log	Complete	`9/10`	Proof tiers and source labels are explicit.
Number source labels	Strong	`9/10`	Main numbers are labeled observed/estimated/assumed. Final scan still required after edits.
AI usage disclosure	Complete	`9/10`	Clear tools, AI role, human decisions, checked items, and limits.
What breaks it	Complete	`9/10`	Written answer and artifact both name bad inputs, missing writeback, no owners, Eric review creep, and cost/model drift.
What stays human	Complete	`9/10`	Pricing, contracts, trust, strategic accounts, people, public claims, and CEO taste/judgment remain human.

Part 1 - First 48 Hours

Prompt line	Current coverage	Score	Notes
Walk through first `48` hours, hour by hour	Strong after artifact update	`8.75/10`	Written answer uses compressed time blocks; operating artifact now has a literal hour-by-hour / window-by-window timeline.
Triage the `23` pending items without knowing the system	Strong	`9/10`	Uses command queue states and source-check routing before deep system knowledge.
Which fire first and why	Strong	`9.25/10`	ClickFlow churn first, stale prospects second, backlog third, Oracle source checks fourth.
What is needed from Eric in first `48` hours	Strong	`9.25/10`	Specific `30`-minute block: priority confirmation, red lines, delegation boundary, `3-5` red decisions, next metric.

Part 2 - Risk Tiering

Prompt line	Current coverage	Score	Notes
Define green/yellow/red with examples	Strong	`9/10`	Agent-specific examples in artifact; concise examples in answer.
What auto-executes and risk tolerance	Strong	`9/10`	Green requires source verified, reversible, known pattern, writeback path, no red flags.
Mistake handling	Strong	`9/10`	Reverse, notify owner, move type to yellow, memory/rule update, repeated-failure rule.
Actual decision tree / criteria	Strong	`9/10`	Value score and risk flags are concrete; red flags override score.

Part 3 - Elon Algorithm Cron Audit

Prompt line	Current coverage	Score	Notes
Keep/kill/improve sample `8` crons	Strong	`9/10`	Kills/pauses Competitor Pricing Monitor and Candidate Outreach Follow-up; improves noisy but valuable crons.
Reasoning framework	Strong	`9/10`	Accepted work, review debt, revenue proximity, source quality, owner clarity, false-positive/false-negative cost, memory feedback.
Avoid keeping everything alive	Strong	`9.25/10`	Explicitly kills at least `2` crons.

Part 4 - Operating Edge

Prompt line	Current coverage	Score	Notes
AI workflow personally automated, before/after metrics	Strong	`9/10`	HG Market Report and HG SEO Commander examples, time savings labeled as estimates.
Evaluate Oracle's `12` quick-win keywords	Strong	`9/10`	Checks source data, URL fit, intent, expected lift, effort, owner, page sensitivity, prior outcomes.
COO/GM in 2028 with AI agents doing `60%` production	Strong	`8.75/10`	Clear human/system boundary; could mention TES revenue mix more often.
Broken/nonexistent operating rhythm	Good	`8.25/10`	Agency drag -> source-data workflow is credible but not a classic inherited-rhythm story.
Enforcing process people resisted	Good	`8/10`	GBP-post CLI example is concrete; operators liked it, so it is more adoption than resistance.
Keep trains running across competing priorities	Strong	`8.75/10`	PASCO + agency + family + software projects + SloaneVault/GTD/Hermes evidence.
MMPI-style assessment	Strong	`9/10`	INTP / `6w5`, Type `1` overlay, Kolbe `9-4-3-2`, explicitly mapped to Eric profile.

What They Evaluate

Evaluation category	Current score	Notes
Operational judgment	`9/10`	Clear priority order, kills, red lines, CEO leverage, churn-first pattern.
AI fluency	`9.25/10`	Uses AI inside operating loops with human boundaries and proof.
Complexity reduction	`8.5/10`	Strong principle, but packet has many internal systems. Keep upload hierarchy lean and use Hermes as backup/public proof, not main clutter.
Source-labeled specificity	`9/10`	Good labels and proof tiers.
Concise communication	`8.5/10`	Main answer is concise; artifact stack needs curation.

What Will Lose - Risk Check

Losing pattern from brief	Do we trigger it?	Read
`48`-hour plan starts with stakeholder meetings	No	We start with command queue, churn, stale prospects, backlog routing.
Risk tiering is conceptual instead of concrete	No	Gate has statuses, scores, risk flags, examples, and failure handling.
Elon audit keeps everything running and adds more	No	We kill/pause at least `2` crons.
"I use ChatGPT for writing SOPs" as AI edge	No	We cite built systems, agent fleet, evals, and operating workflows.
Anything that could apply to any company's operations	Low risk	Single Grain-specific details, agents, ClickFlow/Karrot/TES, Eric profile, and public positioning are used.

What Will Win - Strength Check

| Winning pattern from brief | Current strength | Read |

|---|---|---|
| Triage decisions reveal pattern recognition | `9/10` | Churn and stale pipeline prioritized before generic queue cleanup. |
| Risk tiering implementable Monday | `9/10` | Queue statuses and routing rules are operational. |
| Audit kills at least `2` crons | `9.25/10` | Yes, with reasoning. |
| Evidence of managing automated systems | `9/10` | Hermes, HeavisideOS eval, HG tools. |
| Instinct to close loops, not open workstreams | `9.25/10` | Core thesis is output -> accepted action/rejection/owner/rule update. |

Current Verdict

The packet follows the challenge closely. The two areas to watch are:

1. Do not let the supporting artifacts create complexity after arguing for complexity reduction.
2. If asked live, be ready to explain that the MMPI-style item is the personality/self-assessment work plus Kolbe/Enneagram/MBTI comparison to Eric's profile.

cover-letter

Cover Letter Draft

Dear Single Grain team,

I was not looking for a new role. This one found me.

I follow Eric's content. He is already one of the reference points in my AI agency and company-brain council of experts. When I saw the challenge, I thought it would be fun to test myself against it. Then I read the role more closely. It maps almost exactly to the work I have been doing in my own companies.

I founded Heavyside in 2011 and have run it across SEO, web, digital marketing, and related services. Over the last two years I have rebuilt much of the company around software, AI, and operator-owned workflows. Revenue per FTE has nearly doubled since 2023 while margin improved. My existing businesses now take less than five hours a week to manage because the team and systems run the day-to-day.

The problem in the COO challenge is the problem I keep working on: AI output is cheap. Turning it into accepted work, owner action, revenue, retention, and memory is the hard part. My response lays out how I would do that at Single Grain: protect Eric's time, close loops, kill low-signal automation, and turn ClickFlow, Karrot, and agent workflows into real TES revenue.

More background: <https://csloane.com/numbers>

Chris Sloane

evidence-log

Ops COO 009 - Evidence Log v2

Proof tier reference from `SCORING.md`:

- Tier 0: claims only.
- Tier 1: screenshots.
- Tier 2: demo artifact.
- Tier 3: logs or source records.
- Tier 4: before/after data.
- Tier 5: independent verification.

This evidence log maps the v2 written answer and operating artifact to proof tiers and source labels.

Evidence Table

| Claim | Number source labels | Proof tier | Evidence / source | Shareability / caveat |

|---|---|---|---|---|

| The `ops-coo-009` brief describes `48` crons, `50+` daily messages, `15+` daily review items, `23` pending items, `2` ClickFlow churn customers, `3` stale prospects, `15` inherited Oracle quick-win keywords, and a separate Part 4 evaluation prompt for `12` quick-win keywords from last night. | [Observed from brief] | 3 | `challenges/ops-coo-009/brief.md` | Shareable; public repo brief. |

| The `ops-coo-009` brief says Single Grain's TES/SaaS products are targeting `50%` TES revenue by 2028. | [Observed from brief] | 3 | `challenges/ops-coo-009/brief.md` | This is the strategic reason revenue/FTE should rise through productized/tech-enabled revenue and accepted agent work, not artificial shrinkage. |

| Eric's intended load is `3-5` decisions/day and `15` minutes/day, not `15+` review items/day or `60+` minutes/day. | [Observed from brief] | 3 | `challenges/ops-coo-009/brief.md` | Shareable; public repo brief. |

| The brief's CEO profile says Eric is ENTJ / `8w7`, Kolbe Quick Start `9`, Follow Thru `2`, self-rated weak on follow-through, organization, and workflow consistency, and starts `10` things while finishing `3`. | [Observed from brief] | 3 | `challenges/ops-coo-009/brief.md` | Treat as the central operating clue, not trivia. |

| The GM role is remote and compensation is `\$185K-\$225K DOE`. | [Observed from GM posting] | 3 | `https://careers.singlegrain.com/jobs/567472-general-manager-agency`, verified 2026-05-14. | Recheck before upload because public job pages can change. |

| Single Grain's GM mandate is to increase revenue per employee from `\$180K` to `\$300K-\$400K`. | [Observed from GM posting] | 3 | `https://careers.singlegrain.com/jobs/567472-general-manager-agency`, verified 2026-05-14. | Shareable; public role posting. |

| Single Grain publicly positions the agency around AI-native execution, pipeline/revenue accountability, and scaling results without scaling headcount. | [Observed from public websites] | 3 | `https://www.singlegrain.com/` and `https://www.singlegrain.com/landing/`, verified 2026-05-13. | Use as positioning/context, not as internal performance proof. |

| Single Brain is publicly described as a managed AI agent team inside Salesforce, Gong, GA4, HubSpot, and Slack. | [Observed from public websites] | 3 | `https://www.singlegrain.com/ai-tools/`, verified 2026-05-13. | Public product/positioning page. |

| The standalone Single Brain site frames the architecture as one trusted brain, one commander, and a specialist fleet with isolated client deployments and an SSR trust standard. | [Observed from public websites] | 3 | `https://singlebrain.com/`, verified 2026-05-13. | Public product/positioning page; do not claim actual internal deployment details beyond the page. |

| ClickFlow publicly supports a strategy -> writing -> reporting loop and developer-accessible workflows driven by GSC/GA4, CMS publishing, MCP/REST, and AI visibility data. | [Observed from public websites] | 3 | `https://www.clickflow.com/`, `https://docs.clickflow.com/`, and `https://docs.clickflow.com/capabilities/analytics`, verified 2026-05-13. | Public docs/product pages. |

| Karrot publicly supports LinkedIn ABM personalization, HubSpot pipeline reporting, deal activity timelines, and personalized landing pages. | [Observed from public websites] | 3 | `https://karrot.ai/` and `https://www.singlegrain.com/ai-marketing-tools/`, verified 2026-05-13. | Public product pages. |

| Single Grain's AI Revenue Agents guide emphasizes data quality, bidirectional CRM/marketing automation sync, staged rollout, approval rules, escalation thresholds, recovery protocols, and KPIs. | [Observed from public websites] | 3 | `https://www.singlegrain.com/artificial-intelligence/ai-revenue-agents/`, verified 2026-05-13. | Public thought-leadership page; use for calibration. |

| Single Grain's Company Intelligence guide frames the system as an active intelligence layer across customer behavior, internal ops data, predictive modeling, multi-agent execution, governance, and audit logging. | [Observed from public websites] | 3 | `https://www.singlegrain.com/artificial-intelligence/company-intelligence/`, verified 2026-05-13. | Public thought-leadership page; use for calibration. |

| Eric publicly wrote that Single Grain has rebuilt operations around Claude Code/OpenClaw and saves `90+` hours/month on workflows that previously required dedicated headcount. | [Observed from public commentary] | 3 | `https://newsletter.levelingup.com/p/how-openclaw-claude-code-will-shift-your-business-completely`, verified 2026-05-13. | Public founder commentary, not a verified role fact. |

| Heaviseid started in 2011 and Chris remains founder/operator. | [Observed applicant claim] | 0-1 | Applicant interview notes: `tasks/ops-coo-009-applicant-interview-notes.md`. | Applicant claim; confirm before final upload. |

| Chris spent the last `2` years re-engineering Heaviseid, PVM, and GDM around an intelligence layer, software-driven workflows, and CLI-driven operations. | [Observed applicant claim] | 0-2 | Applicant interview notes; supporting repo/vault artifacts can be summarized if approved. | Applicant claim; share only high-level descriptions unless artifacts are redacted. |

| Chris sustained and grew Heaviseid while holding a full-time PASCO role owning SEO/digital marketing and now ML/AI work. | [Observed applicant claim] | 0-2 | Applicant correction on 2026-05-14; candidate profile and written answer. | Strong talent-leadership and delegation claim. Frame as delegated operating-capacity proof, not divided-attention proof. Applicant-stated unless backed by calendar, PASCO role evidence, operator references, or org chart. |

| Chris's existing businesses take less than `5` hours/week to manage because he hires competent operators, gives ownership, and avoids micromanagement. | [Observed applicant claim] | 0-2 | Applicant correction on 2026-05-14; candidate profile and written answer. | Strong talent-leadership, delegation, and new-role capacity claim. Applicant-stated unless backed by calendar/time logs, operator references, or org chart. |

| Heaviseid has a public operating bench across technical operations, web, SEO, PPC, operations, project flow, client relationship, client success, and reporting/admin support. | [Observed public website / internal ownership map] | 3 | `https://heaviseidgroup.com/about`, `https://heaviseid.digital/about`, checked 2026-05-14; summarized in `heaviseid-org-chart-redacted-summary.md`. | Supports talent leadership, delegation, leadership bench, and founder leverage. Not an FTE or payroll-validation artifact. |

| Heaviseid moved from `\$55.6K` revenue/FTE and `2.6%` net margin in 2023 to `\$77.2K` and `14.2%` in 2024, `\$90.0K` and `12.9%` in 2025, and `\$120.1K` and `17.1%` in 2026 Q1. | [Observed internal P&L / Estimated FTE] | 4 | Native QuickBooks ProfitAndLoss snapshots in Agency Financials plus labor-spend-to-FTE model; summarized in `heaviseid-rpe-profitability-redacted-summary.md`, with internal working calculation available for walkthrough. | Attach the redacted summary only. Do not attach raw confidential P&L exports. |

| Current active-plus-paused Heaviseid MRR is about `\$70.7K/mo`, or `\$848K` annualized, and roughly `\$126K` revenue per paid-team FTE using `6.71` paid-team FTE. | [Observed internal customer/payroll records] | 3-4 | Agency Financials DB query and payroll/FTE calculation summarized in `tasks/ops-coo-009-heaviseid-rpe-calculation.md`. | Useful backup; main written answer uses historical P&L basis instead. |

| HG Market Report implements a lead/report/presentation workflow: business and Google Business Profile capture, AI keyword/service-area setup, SEO audit / SERP / geogrid jobs, derived lost-revenue / quick-win / competitor views, public presentation URL, and report-ready email. | [Observed from private repo docs] | 3 | `hg-market-report` README, `docs/embed/README.md`, `docs/reports/report-architecture.md`, presentation routes, summarized in `heaviseid-sales-delivery-automation-redacted-summary.md`. | Summarize only; do not attach customer reports, private admin screens, env examples, or source excerpts without review. |

| HG Market Report saves roughly `1-2` hours per standard sales audit. | [Estimated applicant claim] | 0-2 | Applicant estimate tied to replaced steps: manual source pulls, local search checks, screenshot/report assembly, presentation creation, and first delivery. | Label as estimate; stronger if later backed by time logs or before/after task records. |

| HG SEO Commander automates large parts of recurring SEO delivery across GA4, Search Console, Google Business Profile, SERP/geogrid tracking, citations, AI articles, GBP posts, bulk-posting approvals, scheduled syncs, BullMQ queues, and CLI/admin tools. | [Observed from private repo docs] | 3 | `hg-seo-commander` README, CLI docs, GBP implementation assessment, scheduled-sync notes, summarized in `heaviseid-sales-delivery-automation-redacted-summary.md`. | Summarize only; strategy and quality review remain human. Stronger with redacted Loom or demo report. |

| HG SEO Commander saves `3-4` hours per active SEO client per month. | [Estimated applicant claim] | 0-2 | Applicant estimate tied to routine data pulls, status checks, report assembly, GBP/content prep, and queue recovery. | Label as estimate; stronger if later backed by time logs. |

| The GBP-post bug/process example shows process enforcement through CLI tooling and adoption by making the operator more capable. | [Observed applicant claim] | 0-2 | Applicant interview notes. | Good narrative proof; not independently verified yet. |

| SloaneVault is a durable memory/source-record system used across projects. | "[Observed internal source]" | 3 | SloaneVault/qmd index, project notes, lessons files, and decision records. | Share high-level structure only, not sensitive vault contents. |

| x-accel provides a proven scoring/gate analogue for AI-generated content before action. | "[Observed internal source]" | 3 | x-accel publish-gate source records and redacted workflow summary. | x-accel is separate from HeavisideOS; summarize unless code excerpts are approved. |

| Decision councils are a reusable judgment mechanism for high-value ambiguous decisions. | "[Observed internal source]" | 2-3 | Decision Council workflow source record and safe/anonymized council artifact. | Use only as supporting proof; do not make every decision a council. |

| HeavisideOS is currently used in practice for SlackBot summaries/query flows around recent client communications and paid ads/SEO/business-performance context. | "[Observed internal source]" | 3 | HeavisideOS Slack connector source records and project lesson. | Do not claim HeavisideOS runs the company end-to-end. |

| HeavisideOS Slack assistant eval improved from overall "4.055/5.000" to "4.467/5.000", later "4.801/5.000", and retrieval improved from "3.255/5.000" to "4.964/5.000" across "117" cases. | "[Observed internal eval]" | 4 | Redacted HeavisideOS Slack assistant eval summary; raw report IDs "20260226-225557", "20260227-015739", and "20260227-025448". |

Optional proof. Submit redacted metrics only; raw reports contain customer names. |

| Chris operates a Hermes/open-agents scheduler fleet with "88" repo-checked jobs, "28" Hermes skills, and "3" managed contexts; live fleet status on 2026-05-14 reported healthy across "3/3" nodes. | "[Observed internal scheduler state]" | 3 | Internal Hermes/open-agents manifests, "sync-hermes.py --verify", and "hermes-fleet.py status --json"; summarized in "hermes-cron-fleet-redacted-summary.md" and "hermes-closed-loop-agent-proof.md". | Bounded cron-fleet operating evidence. "sync-hermes.py --verify" showed no job issues but did show "4" skill content mismatches; do not claim perfect fleet hygiene or every job's business value. |

| Eric/X public commentary emphasizes Single Brain/company-brain workflows, revenue agents, service-as-software, proof-of-work hiring, Slack/Teams adoption, and model/tool-spend governance. | "[Observed captured X corpus]" | 3 | Authenticated x-accel pull of "@ericosiu", refreshed 2026-05-13 through 17:52 UTC. | Public commentary, not role fact; do not overfit current copy to stale posts. |

| The proposed score-to-route gate can reduce Eric's default review burden. | "[Assumed operating design]" | 2 | "tasks/ops-coo-009-operating-artifact-v2.md". | Proposed artifact, not observed Single Grain result. |

| The cron audit kills/pauses "2" sample crons and improves/keeps the rest based on action rate and review debt. | "[Assumed verdict from observed brief data]" | 2-3 | "tasks/ops-coo-009-operating-artifact-v2.md"; sample cron data from public brief. | Verdicts are proposed; action rates are brief facts. |

| MMPI-style self-assessment: founder first; INTP / "6w5" with strong Type "1"; Kolbe "9-4-3-2"; high agency, detail-oriented, systems-minded, ambiguity-tolerant; risk is overbuilding/capability sprawl. | "[Observed applicant self-assessment]" | 1-2 | "[[30 - Personal/37 - Knowledge Graph/people/chris-sloane/personality-and-operating-profile]]", applicant interview notes, candidate profile addendum, and v2 written answer. | Self-assessment, not validated clinical psychometrics; applicant should approve before final upload. |

| AI usage disclosure. | "[Observed workflow]" | 1-2 | "tasks/ops-coo-009-ai-disclosure-v2.md". | Submit final disclosure. |

Claims To Avoid Or Demote

- Do not claim HeavisideOS is a mature company-wide operating system.
- Do not claim x-accel is part of HeavisideOS.
- Do not claim Heaviside is already at Single Grain's "\$300K-\$400K" revenue/FTE target.
- Do not imply Heaviside's FTE reduction was artificial shrinkage; frame it as sensible simplification after overstaffing, supported by improved profitability.
- Do not claim direct agency-scale equivalence to Single Grain's "45"-person organization without a verified source example.
- Do not claim Hermes/open-agents proves every scheduled job is valuable; use it only as bounded evidence of operating a monitored scheduler fleet with pause/error visibility.
- Do not attach raw HeavisideOS eval reports, raw vault notes, private client data, customer screenshots, or raw P&L exports externally without redaction.

Final Recheck

- Reconfirm the GM posting before upload.
- Confirm or replace LinkedIn URL.
- Re-scan attachments for customer names, private client data, secrets, tokens, and raw report tables.
- Confirm applicant approval of self-assessment and operating-experience claims.
- Re-run number-label audit after edits.

heaviside-org-chart-redacted-summary

Heaviside Org Chart - Talent Leadership Summary

Status: draft for applicant review.

Purpose: support the talent-leadership / founder-leverage claim in the Single Grain GM packet without exposing private contact data, customer information, or internal employment details. This is not an FTE, revenue/FTE, or payroll-validation artifact.

Source Basis

Source	What it supports	Notes
`https://heavisidegroup.com/about`	Public named roster and official public titles.	Checked live 2026-05-14. This is the strongest shareable roster source.
`https://heaviside.digital/about`	Public functional structure: Web, Digital Marketing, SEO, and Operations.	Checked live 2026-05-14. It says operations provides sales, customer service, and project-management support.
`[[10 - Business/12 - Areas/Shared/Service Delivery/team-ownership-map]]`	Internal workstream ownership by function and assignee.	Not an HR roster; it is based on observed ClickUp ownership patterns.

Public Roster

The public Heaviside Group About page currently lists the following team:

Person	Public title	Rough function
Chris Sloane	President and Founder	Founder, company leadership, SEO/SEM/web/software direction.
Joe Pfeiffer	Director of Technical Operations	Technical operations, web development, infrastructure-oriented delivery.
Valerie Conaway	Operations Manager	Operations management, platform/reputation/local workflow support.
Dmytro Sylenko	PPC Manager	Paid media, tracking, technical implementation support.
Sean Sudduth	SEO Manager	SEO leadership and local/search delivery.
Jason Teets	Client Success Manager	Client success, sales/marketing experience, customer-facing support.
Mike Moselina	Operations Supervisor	Client project oversight, timelines, deliverables, workflow coordination.
Faith Millares	Client Relationship Specialist	Client relationship support, intake/access/approvals support.
Patrick Reyes	Operations/Process Specialist	Task monitoring, process improvement, deliverable follow-through.
Krystal Foronda	Operations Specialist	Reporting, back-end admin, team support.
Mhikael Divinagracia	Web Designer	Web design and development.
Jhyrfyl Pepito	SEO Specialist	SEO execution and local/search support.

Working Org Chart

This is the practical operating view suggested by the public roster plus internal ownership docs.

```
```text
Chris Sloane - Founder / President
|
|-- Technical Operations / Web
| |-- Joe Pfeiffer - Director of Technical Operations
| |-- Mhikael Divinagracia - Web Designer
| |-- Dmytro Sylenko - PPC Manager / technical implementation support
|
|-- Operations / Project Flow
| |-- Valerie Conaway - Operations Manager
| |-- Mike Moselina - Operations Supervisor
| |-- Patrick Reyes - Operations/Process Specialist
| |-- Krystal Foronda - Operations Specialist
|
|-- Client Success / Account Coordination
| |-- Jason Teets - Client Success Manager
| |-- Faith Millares - Client Relationship Specialist
| |-- Patrick Reyes - process / deliverable follow-through support
|
|-- SEO / Local Visibility
| |-- Sean Sudduth - SEO Manager
| |-- Jhyrfyl Pepito - SEO Specialist
| |-- Valerie Conaway - reputation / GBP workflow support
|
|-- Paid Media / Tracking
| |-- Dmytro Sylenko - PPC Manager
| |-- Faith Millares - access / tracking support
```
```

Functional Pods

| Pod | Public people | Internal workstream evidence |
|---|---|---|
| Founder / strategy / systems | Chris Sloane | Founder remains active, but the agency operates through functional owners rather than daily founder task management. |
| Technical operations and web | Joe Pfeiffer, Mhikael Divinagracia, Dmytro Sylenko | Website staging, WordPress, city/service pages, QA, Volley revisions, technical recovery, tracking implementation. |
| Operations and project flow | Valerie Conaway, Mike Moselina, Patrick Reyes, Krystal Foronda | Package status, agreement upload, billing/rebill dates, client folder completeness, process monitoring, reporting/admin support. |
| Client success and coordination | Jason Teets, Faith Millares, Patrick Reyes | Client intake, access, kickoff, approvals, customer-facing support, project coordination. |
| SEO and local visibility | Sean Sudduth, Jhyrfyl Pepito, Valerie Conaway | SEO units, schema, on-page recommendations, blog/local content, GBP access, reviews, Local Viking, reputation workflows. |
| Paid media and tracking | Dmytro Sylenko, Faith Millares | Google Ads, LSA, conversion tracking, GTM/GA4/GSC, landing-page tracking, ad verification. |
| Social / cold email / ancillary delivery | Krystal Foronda, Valerie Conaway, Joe Pfeiffer, Faith Millares, Shelli Schilke | Social scheduling and cold email setup appear in the ownership map, but public roster coverage is partial. |

Internal-Only Names To Confirm Before Public Use

These names appear in internal docs, continuity notes, or workstream ownership, but should not be added to a public-facing org chart until Chris confirms current status and intended visibility.

| Person | Internal signal | Confirmation needed |
 |---|---|---|
 | Andres Oppenheimer | Account-management and staging-site signals. | Confirm current status and role. |
 | Zachary McCreedy | Operations/development continuity context. | Confirm whether to include in application proof and whether he is public-facing. |
 | Shelli Schilke | Cold email operations / phone-user references. | Confirm current role and whether she belongs in the org chart. |
 | Shannon Conway | HR/admin context. | Confirm whether she owns remote office / HR functions. |
 | Em / Mhikael Divinagracia | Internal notes and public roster use different names. | Confirm whether these are the same person and preferred name. |

Leadership Read

The leadership proof is not that the team exists on a website. The stronger point is that the team is organized into functional ownership areas that can run recurring delivery without day-to-day founder intervention:

- Chris has sustained and grown the agency while holding a full-time PASCO role, which only works if the team can own recurring delivery without founder bottlenecking.
- Chris's existing businesses take less than '5' hours/week to manage, which makes the founder-leverage claim concrete.
- This should be framed as delegated operating-capacity proof, not as a claim that Chris intends to split attention casually. The point is that existing businesses are built to keep running through owners and systems while Chris takes on a new operating role.
- Technical operations and web have named owners.
- SEO, PPC, and local visibility have named functional specialists.
- Operations has a management layer responsible for project flow, task monitoring, process, reporting, and admin support.
- Client relationship / client success has customer-facing coverage.
- Internal ownership docs show work mapped to operators rather than every recurring decision routing back to Chris.

Use In GM Packet

This artifact supports one narrow, useful claim:

> Heavyside has a real operating bench across technical operations, web, SEO, PPC, operations, project flow, client coordination, and reporting. The founder-leverage claim is therefore not just "I use AI"; it is "I hired and organized a team that can own recurring delivery with less than '5' hours/week of founder management, sustained that while also holding a full-time PASCO role, then used software and automation to increase leverage inside that team."

Do not overclaim this as independent proof of exact founder hours. For that, use calendar/time evidence or operator references if needed.

Open Questions

1. Should the application-facing org chart use names plus roles, or role-only pods?
2. Should Jason be shown as Client Success Manager, sales-only, or client success / sales during transition?
3. Should Andres be included as Account Manager?
4. Should Zach be included as operations/development leadership even though he is not on the public About page?
5. Are Em Divinagracia and Mhikael Divinagracia the same person, and which name should be used?

heaviside-rpe-profitability-redacted-summary

Heaviside RPE / Profitability - Redacted Summary

This is a redacted evidence summary for `ops-coo-009`. It is designed to support the GM-role claim without attaching raw QuickBooks reports, customer lists, payroll records, bank data, or private P&L exports.

Why This Is The Primary GM Proof

Single Grain's GM posting centers on moving revenue per employee while maintaining business health. This artifact shows the same operating pattern inside Heaviside: revenue/FTE improved materially while net margin also improved or remained healthy.

This is stronger than a general "I use AI" claim because it ties systems work to the labor math the role is explicitly accountable for.

Source Records

| Evidence | Source reference | Shareability |
|---|---|---|
| Quarterly revenue and net income | Native QuickBooks ProfitAndLoss snapshots synced into Agency Financials on 2026-05-13 | Summarized only; raw P&L exports are confidential. |
| Labor-to-FTE conversion | QuickBooks labor accounts converted through current domestic/non-US cost-per-FTE benchmarks | Summarized method is shareable; raw payroll records are confidential. |
| Current paid-team benchmark | Current payroll allocation: `268.215` hours/week = `6.705` paid-team FTE `[Observed internal payroll records]` | Summarized only. |
| Working calculation | Internal `ops-coo-009` RPE calculation and SloaneVault RPE history note | Summarized only. |

Method

Revenue and net income come from QuickBooks ProfitAndLoss snapshots. FTE is estimated from QuickBooks labor spend by separating domestic and non-US labor pools and converting those pools through current cost/FTE benchmarks.

The method is imperfect but directionally useful: it uses observed finance records and a stated FTE conversion method rather than headcount guesses or bank-deposit cash flow.

Year-Level Summary

| Year / period | P&L revenue | Avg FTE-equivalent | Revenue/FTE read | Net income | Net margin |
|---|---|---|---|---|---|
| 2022 | `\$589,535` `[Observed internal P&L]` | `15.80` `[Estimated FTE]` | `\$37,306/FTE` `[Observed internal P&L / Estimated FTE]` | `\$3,848` `[Observed internal P&L]` | `0.7%` `[Observed internal P&L]` |
| 2023 | `\$812,398` `[Observed internal P&L]` | `14.61` `[Estimated FTE]` | `\$55,606/FTE` `[Observed internal P&L / Estimated FTE]` | `\$21,201` `[Observed internal P&L]` | `2.6%` `[Observed internal P&L]` |
| 2024 | `\$783,177` `[Observed internal P&L]` | `10.15` `[Estimated FTE]` | `\$77,198/FTE` `[Observed internal P&L / Estimated FTE]` | `\$111,188` `[Observed internal P&L]` | `14.2%` `[Observed internal P&L]` |
| 2025 | `\$632,867` `[Observed internal P&L]` | `7.03` `[Estimated FTE]` | `\$90,024/FTE` `[Observed internal P&L / Estimated FTE]` | `\$81,546` `[Observed internal P&L]` | `12.9%` `[Observed internal P&L]` |
| 2026 Q1 annualized | `\$185,710` quarterly `[Observed internal P&L]` | `6.27` `[Estimated FTE]` | `\$120,094/FTE` `[Observed internal P&L / Estimated FTE]` | `\$31,712` quarterly `[Observed internal P&L]` | `17.1%` `[Observed internal P&L]` |

Strategic Read

The cleanest story starts in 2024, when AI/system-integration work began lining up with the operating model.

| Period | Revenue/FTE | Net margin | Read |
|---|---|---|---|
| 2023 | `\$55,606/FTE` `[Observed internal P&L / Estimated FTE]` | `2.6%` `[Observed internal P&L]` | Labor-heavy baseline. |
| 2024 | `\$77,198/FTE` `[Observed internal P&L / Estimated FTE]` | `14.2%` `[Observed internal P&L]` | AI/systems integration and simplification start showing up. |
| 2025 | `\$90,024/FTE` `[Observed internal P&L / Estimated FTE]` | `12.9%` `[Observed internal P&L]` | Efficiency holds with healthy profitability. |
| 2026 Q1 | `\$120,094/FTE` `[Observed internal P&L / Estimated FTE]` | `17.1%` `[Observed internal P&L]` | Trajectory continues. |

This supports the role-specific claim: higher revenue per paid-team FTE and better profitability moved together. It should be framed as sensible simplification after overstaffing plus systems/automation, not artificial shrinkage.

Submission Use

Recommended phrasing:

> On a QuickBooks P&L basis, with FTE estimated from labor spend, Heaviside moved from `\$55.6K` revenue/FTE and `2.6%` net margin in 2023 to `\$77.2K` and `14.2%` in 2024, `\$90.0K` and `12.9%` in 2025, and `\$120.1K` and `17.1%` in 2026 Q1 `[Observed internal P&L / Estimated FTE]`. That is not Single Grain's target yet. It is proof that I understand the metric and can move it while business health improves.

Caveats

- Do not attach raw QuickBooks exports externally without redaction.
- FTE is estimated from labor spend, not a perfect person-by-person historical roster.
- 2026 Q2TD exists internally but should not be overused because partial-quarter timing is lumpy.
- This does not claim Heaviside is already at Single Grain's `\$300K-\$400K` revenue/employee target.

heaviside-sales-delivery-automation-redacted-summary

Heaviside Sales / Delivery Automation - Redacted Summary

This is a redacted evidence summary for `ops-coo-009`. It supports the operating-edge claim that Chris has built software systems that remove manual sales and delivery work inside his own agency companies.

Do not attach raw customer reports, customer screenshots, live production admin pages, API keys, `.env` files, or private client examples.

What This Proves

The role is looking for a builder of operational leverage, not a traditional agency coordinator. These systems show that pattern:

- turn manual report/deck work into generated customer-facing artifacts,
- turn delivery tasks into queueable/reviewable workflows,
- keep humans on review, judgment, customer communication, and exceptions,
- move repeatable work into software, CLI, queues, and integrations.

Artifact 1: HG Market Report

What It Is

HG Market Report is a local-business market intelligence platform. The repo supports SEO audits, SERP analysis, local geogrid visibility, competitor analysis, review monitoring, AI visibility, Google Ads analysis, lead scoring, email campaigns, and a white-label lead capture widget.

The most relevant workflow is the lead/report/presentation path: the widget captures business and Google Business Profile context, AI-generates keywords and service area, creates lead and business records, queues SEO audit / SERP / geogrid reports, triggers derived lost-revenue / quick-win / competitor views, creates a public presentation URL, tracks presentation views, and sends a report-ready email.

Source Records

| Evidence | Source reference | Shareability |
|---|---|---|
| Product repo | `hg-market-report` repository | Summarize only unless code excerpts are reviewed for private identifiers. |
| Project note | SloaneVault `HG Market Report` project note | Summarize only. |
| Production surface | `hgmarketreport.com` | Public surface can be named; do not expose private admin/customer data. |
| Repo README / docs | README, `docs/embed/README.md`, `docs/reports/report-architecture.md`, presentation routes, worker/runbook docs | Safe to summarize after redacting env examples and private deployment details. |

Operating Claim Supported

Repo-supported claim: HG Market Report replaces the manual prospect-audit loop with a software path from captured business context to queued reports, derived analysis, public presentation URL, and report-ready email.

Applicant estimate: it conservatively saves `1-2` hours per standard sales audit `[Estimated applicant claim]` by removing manual source pulling, local search checks, screenshot/report assembly, presentation creation, and first delivery. More complex multi-report audits may save more, but the public claim should stay at `1-2` hours unless backed by time logs.

Proof Tier Read

- Product/repo/source records: Tier 3 if source records are reviewed live.
- Demo artifact: Tier 2 if a redacted report or Loom walkthrough is attached.
- Time-savings estimate: Tier 0-2 unless backed by task logs or before/after timestamps.

Artifact 2: HG SEO Commander

What It Is

HG SEO Commander is a multi-tenant local SEO automation platform. It integrates GA4, Google Search Console, Google Business Profile, SERP/geogrid tracking, citation monitoring, AI content generation, queue workers, scheduled syncs, CLI tools, and operator-review workflows.

The most relevant workflow is monthly SEO delivery: routine data pulls, rank/local visibility checks, citation context, Google Business Profile posts, content generation, scheduled sync, queue handling, and admin/CLI recovery are moved into software. Strategy, customer judgment, and quality review remain human.

Source Records

| Evidence | Source reference | Shareability |
|---|---|---|
| Product repo | `hg-seo-commander` repository | Summarize only unless code excerpts are reviewed for private identifiers. |
| Project note | SloaneVault `HG SEO Commander` project note | Summarize only. |
| Production surface | `hgseocommander.com` | Public surface can be named; do not expose private admin/customer data. |
| Repo README / docs | README, CLI docs, GBP implementation assessment, scheduled-sync notes, API docs, queue docs | Safe to summarize after redacting env examples, customer IDs, and account IDs. |

Operating Claim Supported

Repo-supported claim: HG SEO Commander automates large parts of the recurring SEO operating workflow across GA4, Search Console, Google Business Profile, SERP/geogrid, citations, AI articles, GBP posts, bulk-posting approvals, scheduled syncs, BullMQ queues, and CLI/admin tools.

Applicant estimate: it conservatively removes `3-4` hours per active SEO client per month `[Estimated applicant claim]` from routine data pulls, status checks, report assembly, GBP/content prep, and queue recovery. The stronger proof is not the exact hours yet; it is that the platform has a real production architecture for delivery work: Google integrations, queue workers, approval flow, CLI operations, content generation, publishing controls, and recoverability.

Proof Tier Read

- Product/repo/source records: Tier 3 if source records are reviewed live.
- Demo artifact: Tier 2 if a redacted walkthrough or sanitized screen capture is attached.

- Time-savings estimate: Tier 0-2 unless backed by time logs.

How To Use In The Submission

Do not overload the main written answer. Use this as backup for the Operating Edge section:

> I have built a lot of internal tools; two clean examples are HG Market Report and HG SEO Commander. HG Market Report turns a prospect audit into a workflow: captured business context, queued reports, derived analysis, public presentation URL, and report-ready delivery. HG SEO Commander does the same thing on recurring SEO delivery across analytics, search visibility, GBP, citations, content, queues, and CLI/admin recovery. The important point is not that the tools are perfect. It is that I repeatedly move agency work from manual coordination into source-backed, reviewable software workflows.

Caveats

- Do not claim these two tools are the only systems Chris has built.
- Do not claim exact time savings as observed unless backed by time logs.
- Do not expose customer reports, customer names, account IDs, API keys, OAuth credentials, or raw admin screens externally.
- Frame HG SEO Commander as real delivery automation while still distinguishing it from early-stage HeavisideOS.

heavisideos-slack-eval-redacted-summary

HeavisideOS Slack Assistant Eval - Redacted Summary

This is a redacted evidence summary for `ops-coo-009`. It is designed to be shareable without raw customer names or raw report tables.

Why This Is The Primary Part 4 Proof

This is the cleanest operating-edge proof because it shows a real AI workflow with:

- a repeatable eval runner,
- production API benchmark reports,
- explicit scoring dimensions,
- gate status,
- before/after results,
- code changes tied to the improvement,
- and a documented learning loop.

It should be used as proof of measured AI assistant improvement, retrieval discipline, and operating-loop behavior. It should not be used to claim HeavisideOS is a fully mature company operating system.

Source Records

| Evidence | Source reference | Shareability |
|---|---|---|
| Eval runbook | HeavisideOS Slack Assistant Usefulness Eval Program | Available as summarized/redacted source record. |
| Baseline report | `slack-usefulness/20260226-225557/report.md` | Raw report contains customer names; not attached. |
| Final pass report | `slack-usefulness/20260227-015739/report.md` | Raw report contains customer names; not attached. |
| Later structured-analysis pass | `slack-usefulness/20260227-025448/report.md` | Raw report contains customer names; not attached. |
| Source commit | HeavisideOS commit `1dd49b8` | Shareable if repo/source access is appropriate. |
| Eval runner | `scripts/assistant-evals/run-slack-assistant-eval.ts` | Shareable if repo/source access is appropriate. |
| Eval config | `config/evals/slack-assistant/usefulness-suite.json` | Shareable if tenant data is not embedded. |

Redacted Metrics Table

| Metric | Baseline | Final pass | Delta |
|---|---|---|---|
| Active mapped customers | 39 `[Observed internal eval]` | 39 `[Observed internal eval]` | no change |
| Questions per tenant | 3 `[Observed internal eval]` | 3 `[Observed internal eval]` | no change |
| Total eval cases | 117 `[Observed internal eval]` | 117 `[Observed internal eval]` | no change |
| Overall score | 4.055 / 5.000 `[Observed internal eval]` | 4.467 / 5.000 `[Observed internal eval]` | +0.412 `[Observed internal eval]` |
| Root-cause score | 5.000 / 5.000 `[Observed internal eval]` | 5.000 / 5.000 `[Observed internal eval]` | 0.000 `[Observed internal eval]` |
| Evidence score | 4.872 `[Observed internal eval]` | 4.859 `[Observed internal eval]` | -0.013 `[Observed internal eval]` |
| Actionability score | 3.705 `[Observed internal eval]` | 3.991 `[Observed internal eval]` | +0.286 `[Observed internal eval]` |
| Retrieval score | 3.255 `[Observed internal eval]` | 4.964 `[Observed internal eval]` | +1.709 `[Observed internal eval]` |
| Formatting score | 2.987 `[Observed internal eval]` | 3.000 `[Observed internal eval]` | +0.013 `[Observed internal eval]` |
| Gate status | `review\_required` `[Observed internal eval]` | `pass` `[Observed internal eval]` | gate cleared |
| API responses passing | not recorded in baseline table | 117 `[Observed internal eval]` | final pass complete |
| Retrieval non-empty pass | baseline had 3 failures `[Observed internal eval]` | 117 pass / 0 fail `[Observed internal eval]` | tenant retrieval gap cleared |
| Citation presence pass | baseline had 3 warnings `[Observed internal eval]` | 117 pass `[Observed internal eval]` | citation warnings cleared |

Later structured-analysis pass (`20260227-025448`): overall score reached 4.801 / 5.000 `[Observed internal eval]`, actionability reached 5.000 / 5.000 `[Observed internal eval]`, retrieval remained 4.964 / 5.000 `[Observed internal eval]`, and the gate remained `pass` `[Observed internal eval]`.

What Changed

- Implemented a Slack assistant usefulness eval runner and config.
- Added service-focused golden cases across active PVM/GDM customers.
- Implemented baseline service architecture: catalog defaults, tenant overrides, and service fact overrides.
- Added Slack `SERVICE` fact-update path with tenant service sync on approval.
- Remediated tenant-specific communications retrieval gaps with targeted ClickUp/Missive sync.
- Added a safe Slack test lane so live Slack-path validation can run without posting into customer channels by default.

Submission Use

Use this as the primary answer to: "What's one operational workflow you've personally automated or augmented with AI? Give us the before/after with real metrics."

Recommended phrasing:

> I built an eval loop for a Slack assistant that answers client-status, service, and root-cause questions from customer context. The first production benchmark across 39 active mapped customers `[Observed internal eval]` and 117 cases `[Observed internal eval]` showed an overall score of 4.055/5.000 `[Observed internal eval]` and retrieval of 3.255/5.000 `[Observed internal eval]`, with the gate requiring review. After targeted retrieval/service-state fixes and tenant sync remediation, the full pass reached 4.467/5.000 overall `[Observed internal eval]`, 4.964/5.000 retrieval `[Observed internal eval]`, and a passing gate `[Observed internal eval]`. The lesson was that aggregate quality can hide a binding category failure, so I added category gates, source checks, and repeatable regression reports.

Caveats

- Raw reports include customer names; external packet should use this redacted metrics table.
- The workflow proves measured assistant improvement and retrieval/action-loop discipline.
- It does not prove HeavisideOS is fully mature as a company-wide operating system.
- It maps to Single Grain because the problem is structurally similar: Slack/data-connected AI outputs must be measured by usefulness, retrieval, accepted work, and closed-loop improvement.

hermes-closed-loop-agent-proof

Hermes Closed-Loop Agent Proof

Status: support artifact / public-page source material. Do not attach raw manifests unless requested.

Purpose: connect Chris's Hermes/open-agents experience to the `ops-coo-009` problem: AI systems generating outputs that must be converted into outcomes.

Source Basis

| Source | What it supports |
|---|---|
| `/home/cvsloane/dev/open-agents/ARCHITECTURE.md` | Skills -> Agents -> Hermes Jobs maturity model; fleet topology; data flow; Discord routing; Redis/run state. |
| `/home/cvsloane/dev/open-agents/hermes/jobs/open-agents.json` | HQ scheduled jobs for infra, queue health, backups, project status, security, incidents, research, content, inbox, fleet, X/content, and business summaries. |
| `/home/cvsloane/dev/open-agents/hermes/jobs/cold-email.json` | Cold-email operations jobs for health, deliverability, copy audit, experiment management, reply ingest/triage, report fulfillment, GHSL sync, and pod seeding/top-up. |
| `/home/cvsloane/dev/open-agents/hermes/jobs/agent-commander.json` | Agent Commander watchdog, governance digest, and daily briefing jobs. |
| Hermes live checks on 2026-05-14 | `sync-hermes.py --verify` checked `88` jobs, `28` skills, `3` contexts with no job issues; `hermes-fleet.py status --json` reported the fleet healthy across `3/3` nodes. |

What Hermes Does

Hermes is the production scheduling and routing layer for open-agents. The operating pattern is:

```
``text
scheduled trigger
-> prompt / skill / context
-> agent or script execution
-> local output and trace
-> routing decision
-> Discord / local surface / alert
-> follow-up, pause, retry, or memory/process update
````
```

That is the same shape as the Single Grain problem: outputs are only useful when they get routed, reviewed, accepted, rejected, escalated, or written back into an operating system.

### ## Current Workflow Categories

| Category | Example jobs | Outcome loop |  
|---|---|---|  
| Infrastructure and queue health | Infra Health, Queue Health, DB Backup, DB Anomaly, Disk Usage, Certificate Expiry, DNS Verification | Outputs route to local/monitoring surfaces; warnings/errors trigger investigation instead of passive reporting. |  
| Fleet operations | Fleet Summary, Fleet Alert Rollup, Apps VPS Summary, HomeLinux Summary | Worker outputs are relayed to the right surface; cross-node health becomes an operating status, not separate logs. |  
| Incident and watchdog loops | Incident Responder, Incident TTL Sweep, Watchdog, Control Plane Burn-In | Jobs produce retry/alert/status loops with pause/error visibility. |  
| Research and content | Research: Dev Tools & AI Agents, Marketing & Ads, Cold Email & Deliverability, Client Industry Intel, Research Content Drafter | Research outputs become candidate content drafts or approval items rather than loose notes. |  
| Business and marketing ops | Meta Ads Daily Briefing, Content Strategy, Engagement Digest, Weekly X Growth Report, X Watchlist, inbound/competitor monitoring | Agent outputs are routed into operating decisions around ads, content, X, and market signals. |  
| Inbox / executive assistant | EA Inbox Processor, M365/GWS inbox triage, follow-up sweep, meeting prep, day brief | Information gets triaged into briefs, follow-ups, or next actions. |  
| Cold email operations | Health check, deliverability watchdog, copy audit, experiment manager, reply ingest/triage, daily report, GHSL sync, pod seeding/top-up | Cold-email output can move into replies, triage, reporting, GHSL updates, and campaign/copy operations. |  
| Agent Commander | Watchdog, governance digest, daily briefing | Agent run state becomes a governance surface rather than invisible background activity. |

### ## Evidence Boundaries

- The packet should continue to use Hermes as bounded proof of operating a monitored scheduled-agent / cron fleet.
- Do not claim every Hermes job is high-value.
- Do not claim cold email is clean current performance proof unless live credentials and launch state are verified.
- The useful proof is operating maturity: jobs have manifests, schedules, skills, routing, output surfaces, live status, pause/disabled states, retries, and alerts.

### ## Application Use

Best concise phrasing:

> I operate Hermes/open-agents as a closed-loop scheduler and agent layer. It handles research, content drafts, infra/queue health, project status, inbox/briefing flows, fleet monitoring, and cold-email operations. The point is not more AI output. The point is that outputs have routes: alert, owner action, approval draft, report, CRM/GHL sync, pause, retry, or rule update.

Use this in interviews, cover letter context, or as supporting context behind `csloane.com/numbers`. Keep the main challenge answer focused on Single Grain's operating system.

# hermes-cron-fleet-redacted-summary

# Hermes / Open-Agents Scheduler Fleet - Redacted Summary

Purpose: support the `ops-coo-009` cron-audit section with bounded evidence that Chris operates a real scheduled-agent / cron fleet, not just a proposed cron framework.

## Snapshot

Checked: 2026-05-14.

Source records:

- Internal repo manifest: `open-agents/hermes/jobs/open-agents.json`
- Internal repo manifest: `open-agents/hermes/jobs/cold-email.json`
- Internal repo manifest: `open-agents/hermes/jobs/agent-commander.json`
- Local Hermes scheduler state via `sync-hermes.py --verify` and `hermes-fleet.py status --json`
- Internal monitor guide: `open-agents-monitor/DETAILED\_GUIDE.md`

## Observed Counts

| Source                               | Count / | status |
|--------------------------------------|---------|--------|
| --- ---                              |         |        |
| Repo-verified Hermes jobs            | `88`    |        |
| Hermes skills checked                | `28`    |        |
| Managed contexts checked             | `3`     |        |
| Job-manifest issues                  | `0`     |        |
| Skill content mismatches             | `4`     |        |
| Fleet nodes healthy                  | `3/3`   |        |
| HQ jobs healthy/scheduled            | `61/61` |        |
| Cold-email jobs intentionally paused | `19`    |        |
| Worker jobs reporting `ok`           | `6/6`   |        |

## What This Proves

- Chris operates a real scheduled-agent / cron fleet with more than `80` checked jobs.
- The fleet has explicit pause/disable state, last-run status, next-run timing, routing, delivery surfaces, stale thresholds, and monitoring checks.
- The operating posture is not "keep every cron on." The cold-email manifest currently has `19` intentionally paused jobs pending launch-state verification, which is consistent with the packet's kill/pause/tighten philosophy.
- The system exposes drift instead of hiding it: verification showed `4` skill content mismatches even though job manifests had no issues.

## Caveats

- This is Tier 3 internal source-record evidence, not independent verification.
- This does not prove every scheduled job is high-value.
- This does not prove enterprise-scale Single Grain operations experience.
- Do not attach raw scheduler state externally without reviewing for private routing/channel metadata.

## Safe Submission Wording

> I also operate a Hermes/open-agents scheduler fleet. On 2026-05-14, repo verification checked `88` jobs, `28` skills, and `3` managed contexts with no job-manifest issues, and live fleet status reported healthy across `3/3` nodes "[Observed internal scheduler state]". I would treat that as proof of cron-fleet operating exposure, not proof that every cron deserves to exist. The useful lesson is the same as my Single Grain plan: crons need owners, monitoring, pause/kill states, error visibility, and accepted-work feedback.

# job-post-line-scorecard

# GM Job Post Line-By-Line Scorecard

Status: working review document for application QA.

Source: Single Grain General Manager (Agency) job post, rechecked 2026-05-14: `https://careers.singlegrain.com/jobs/567472-general-manager-agency`.

Purpose: score the application packet against the exact "not a fit" and "strong candidates" lines, not just the general role vibe.

## Exclusion Criteria

These are the lines where the job post says the role is not a fit if experience is primarily the listed pattern.

| Job-post exclusion line | Current strength against exclusion | Score | Evidence in packet | Gap / next proof |  
|---|---|---|---|---|  
| Managing teams without improving systems | Very strong against exclusion | `9.5/10` | HG Market Report, HG SEO Commander, HeavisideOS Slack eval, Hermes/open-agents, RPE/margin history. | Add one concise talent-management example if space exists, but do not bloat. |  
| Coordinating projects rather than building infrastructure | Very strong against exclusion | `9.5/10` | FDE framing, source-data -> generated deliverable -> operator review -> customer send/publish -> outcome log, operating artifact, Hermes proof. | Keep infrastructure examples concrete and avoid internal-tool sprawl. |  
| Relying on founder involvement to resolve operational issues | Strong against exclusion | `9/10` | Existing businesses take less than `5` hours/week to manage; public org chart and ownership map; PASCO/full-time role while agency sustained/grown. | Stronger with calendar/time log or operator reference. |  
| Growing headcount without improving efficiency | Very strong against exclusion | `9/10` | 2023 -> 2026 RPE and margin history; sensible simplification after overstaffing; not artificial shrinkage. | Explain hiring discipline if asked; raw finance should remain redacted. |

Overall exclusion-risk read: low if framed as operational leverage and delegated capacity. Main watchout is the existing-businesses point; it must read as proof that the businesses are delegated/system-run, not that Chris intends to split attention casually.

## Strong-Candidate Lines

| Job-post strong-candidate line | Current fit | Score | Evidence in packet | Gap / next proof |  
|---|---|---|---|---|  
| Operated or scaled a performance marketing agency or service business | Strong | `8.5/10` | Founder/operator of Heaviside since 2011; SEO, web design, digital marketing, related services; PVM/GDM vertical niches; sales/delivery systems. | Smaller than Single Grain; clarify founder/operator scale honestly. |  
| Built operational systems that significantly improved productivity | Very strong | `9.25/10` | HG Market Report saves estimated `1-2` hours/report; HG SEO Commander saves estimated `3-4` hours/client/month; HeavisideOS eval improvement; Hermes fleet. | Time-savings estimates should remain labeled unless backed by time logs. |  
| Demonstrated strong financial discipline | Strong | `8.75/10` | P&L-based RPE/net margin history; improved productivity and margin together; avoids artificial shrinkage. | Could add one hiring/cost-governance story if needed. |  
| Built leadership teams capable of operating independently | Strong | `8.75/10` | Less than `5` hours/week founder management; PASCO/full-time role; public roster; ownership map; org chart proof. | Third-party/operator references would lift this. |  
| Used AI or automation to increase operational leverage | Very strong | `9.5/10` | HG Market Report, HG SEO Commander, HeavisideOS, Hermes/open-agents, x-accel gate, SloaneVault, councils. | Keep proof hierarchy simple to avoid complexity penalty. |

## Candidate Profile Lines

| Job-post profile line | Current fit | Score | Evidence |  
|---|---|---|---|  
| Former agency GMs or COOs | Partial by title, strong by function | `7.75/10` | Founder/operator has done GM/COO work inside own agency, but not titled as outside agency GM. |  
| Operators from performance marketing agencies | Strong | `8.5/10` | Heaviside/PVM/GDM delivery, paid media/SEO/web ops, sales team and agency systems. |  
| Founders who scaled service businesses | Very strong | `9/10` | Heaviside since 2011, vertical niches, agency runs with limited founder management. |  
| Operational discipline | Strong | `8.75/10` | GTD/ToDoist, SloaneVault, command queue, gates, scorecards, cadence, memory. |  
| Strong hiring judgment | Strong but mostly applicant-stated | `8.5/10` | Agency operating bench, under `5` hours/week founder management, PASCO parallel role. |  
| Financial rigor | Strong | `8.75/10` | P&L/RPE/margin model, not artificial shrinkage, revenue/FTE focus. |  
| Systems thinking | Very strong | `9.5/10` | Entire packet is built around systems, source records, writeback, dashboards, and memory. |  
| AI-forward mindset | Very strong | `9.5/10` | AI workflows, agent fleet, AI usage disclosure, human boundaries, model/tool-spend awareness. |

## Core Responsibility Fit

| Responsibility | Fit score | Read |  
|---|---|---|  
| Operational leadership | `8.75/10` | Strong plan and artifact; needs live SG economics/team access to calibrate. |  
| Financial discipline | `8.75/10` | Strong RPE/margin proof; can improve with hiring-decision story. |  
| Talent leadership | `8.75/10` | Stronger after adding less-than-`5`-hours/week proof; still mostly self-reported. |  
| Founder leverage | `9.5/10` | One of the strongest matches. The plan explicitly reduces Eric's review load. |  
| Systems and AI infrastructure | `9.5/10` | Very strong and central to the packet. |

## Current Verdict

The application is aligned with the job post's real screen. It does not read like a traditional agency operator, project coordinator, founder-dependent operator, or headcount-growth operator. It reads like a founder/operator who builds leverage systems, has measurable RPE/margin proof, and understands the CEO-leverage problem.

Next improvement should be selective: one concrete talent/accountability story or one operator reference would improve the only remaining weaker area. Do not add more tool names unless they support the "outputs -> outcomes" thesis.

# live-role-check

# Ops COO 009 - Live Role Check v2

Checked: 2026-05-14.

Source: <https://careers.singlegrain.com/jobs/567472-general-manager-agency>

## Confirmed Posting Facts

| Field | Current value |  
|---|---|  
| Role | General Manager (Agency) |  
| Department | Admin |  
| Location | Remote |  
| Remote status | Fully Remote |  
| Compensation | `\$185,000 - \$225,000 DOE` |  
| Required challenge | COO Operations Challenge / `ops-coo-009` |  
| Submission format | PDF or Markdown, max `4` pages on the job post; the challenge brief sets the written answer cap at `2` pages. |

## Role Mandate Confirmed

The posting still frames the role as a builder of operational leverage, not a traditional agency operator.

Key fit points for the v2 packet:

- Increase revenue per employee from `\$180K` to `\$300K-\$400K`.
- Improve CEO leverage.
- Build scalable operating infrastructure.
- Maintain business health and avoid artificial shrinkage.
- Use AI-enabled workflows where humans focus on judgment and systems handle coordination, analysis, and repetitive work.

# operating-artifact

# Ops COO 009 - Operating Artifact v2

Purpose: turn Single Grain's agent output into measurable leverage: accepted work, owner action, memory updates, lower CEO review load, and a path toward higher revenue per employee.

Number-source rule: challenge facts are labeled "[Observed from brief]"; proposed thresholds, cadence, and operating design are labeled "[Assumed]" until calibrated against Single Grain's live queue, CRM, analytics, finance, and team capacity.

Public-site calibration: Single Grain publicly sells AI-native execution inside the customer's operating stack, not disconnected AI dashboards. The implementation should assume source systems like Slack, HubSpot/Salesforce, Gong, GA4, GSC, CMS, LinkedIn, ClickFlow, Karrot, and Single Brain. The operating standard should match the public Single Brain promise: one trusted interface, a commander/fleet behind it, clear status, client/data isolation, and Security/Stability/Reliability "[Observed from public websites]".

## ## 1. Revenue/FTE Operating Loop

```
```text
source data
-> agent output
-> source check
-> value score + risk flags
-> green / yellow / red route
-> owner execution
-> system-of-record writeback
-> accepted-work log
-> outcome + cost record
-> memory/rule update
-> revenue/FTE review
```
```

Core rule: an AI output does not count until it becomes an accepted action, explicit rejection, owner assignment, or rule/memory update.

Single Grain language version: strategy audit and quick wins -> growth roadmap -> rapid execution -> scale what works "[Observed from public websites]".

Primary operating goal: reduce Eric from `60+` minutes/day of approvals "[Observed from brief]" to `15` minutes/day "[Observed from brief]" and `3-5` real decisions/day "[Observed from brief]", while tying AI workflows to retention, pipeline, margin, productivity, and revenue/FTE.

## ## 2. CEO Operating Contract

The brief gives the operating clue: Eric is ENTJ / `8w7`, Kolbe Quick Start `9`, Follow Thru `2`, self-rated low on follow-through, organization, and workflow consistency, and "starts `10` things, finishes `3`" "[Observed from brief]". Do not make the system depend on his follow-through.

| Eric should do | GM / system should do |  
|---|---|  
| Set strategic direction, taste, standards, partnerships, and red-line decisions. | Convert starts into owned work, status, dates, decision logs, and outcomes. |  
| Use his Quick Start energy to create options and find leverage. | Keep a `10` starts -> `10` closed loops` ledger "[Assumed]": shipped, killed, delegated, parked, or escalated. |  
| Spend `15` minutes/day "[Observed from brief]" on true red decisions. | Protect him from yellow/green review debt and summarize only exceptions. |  
| Jump to the next big idea without carrying the process in his head. | Maintain memory, owners, writebacks, and follow-through outside the CEO's attention. |

Applicant fit: Chris's profile is INTP / `6w5` with strong Type `1`, Kolbe `9-4-3-2` "[Observed self-assessment]". Practical use: source-checking, preparedness, standards, and error-correction as the counterweight to high-Quick-Start CEO energy.

## ## 3. First 48-Hour Command Queue

| Status | Meaning | SLA |  
|---|---|---|  
| `dead` | Duplicate, stale, irrelevant, or no longer recoverable. | Same-day close "[Assumed]" |  
| `needs\_source\_check` | Potentially useful, but data/source is unverified. | `4` business hours "[Assumed]" |  
| `green\_auto` | Low-risk, reversible, source-verified, known pattern. | Same day "[Assumed]" |  
| `yellow\_owner` | Needs GM/domain-owner review, not Eric. | `4` business hours "[Assumed]" |  
| `yellow\_council` | High-value ambiguity needing expert judgment. | `24` hours "[Assumed]" |  
| `red\_eric` | Strategic, irreversible, public, pricing, legal, people, or key-relationship risk. | Eric's `30`-minute block "[Observed from brief]" |  
| `executed` | Action taken and owner assigned. | Log within `24` hours "[Assumed]" |  
| `learning\_logged` | Outcome changed a rule, prompt, threshold, or SOP. | Weekly sweep "[Assumed]" |

## ## 3A. First 48-Hour Takeover Timeline

The written answer compresses this for page fit. The operating version is more literal:

| Window | Action | Outcome |  
|---|---|---|  
| Hour `0-1` "[Assumed]" | Snapshot Mission Control, review queue, ClickFlow churn flags, expired outreach, CRM, product usage, Slack, GSC/GA4, and source links. | One source-backed command queue. |  
| Hour `1-2` "[Assumed]" | Classify all `23` pending items "[Observed from brief]" into `dead`, `needs\_source\_check`, `green\_auto`, `yellow\_owner`, `yellow\_council`, or `red\_eric`. | No unknown/unclassified backlog. |  
| Hour `2-3` "[Assumed]" | Pull ClickFlow churn context for the `2` lost customers "[Observed from brief]": usage, support, contract value, last owner action, save path. | Churn recovery packets ready. |  
| Hour `3-5` "[Assumed]" | Assign churn recovery owner and decide whether any item requires Eric because of pricing, strategic relationship, or customer-trust risk. | Churn routed before backlog cleanup. |  
| Hour `5-6` "[Assumed]" | Review `3` stale prospects "[Observed from brief]": CRM stage, last reply, account fit, trigger, relationship risk, owner. | Re-entry or dead-item decision per prospect. |  
| Hour `6-8` "[Assumed]" | Clear routine greens, mark dead items, and move source-uncertain items to source check. | Queue size collapses without using Eric as reviewer. |  
| Hour `8-10` "[Assumed]" | Prepare Eric's `3-5` red decisions "[Observed from brief]" plus red-line/delegation asks. | Eric gets decision packet, not discovery. |  
| Hour `10-24` "[Assumed]" | Execute green/yellow owner actions, write back decisions, and schedule day-two Oracle/Shawn work. | Day one ends with no unowned yellow/red item. |  
| Day `2`, morning "[Assumed]" | Meet Eric for `30` minutes "[Observed from brief]"; meet Shawn to identify agent failure modes and dashboard/writeback gaps. | Delegation boundary plus tech constraint map. |



| Day `2`, afternoon `[Assumed]` | Source-check Oracle's `15` inherited quick-win keywords `[Observed from brief]`; route winners into owner actions and losers into memory. | SEO outputs become accepted work or explicit rejections. |  
| Day `2`, close `[Assumed]` | Publish accepted-work, rejected-work, red-decision, owner, and rule-update summary. | The system has started converting output into outcomes. |

## 4. Day-One Fire Order

| Fire / item group | Initial route | Why | First action | Eric needed? |  
|---|---|---|---|  
| `2` ClickFlow churn customers `[Observed from brief]` | `yellow\_owner` -> possible `red\_eric` | Active TES/SaaS revenue loss outranks backlog cleanup. | Pull product usage, GSC/GA4 where relevant, support history, contract value, churn signal, save path, and owner. | Only for pricing, strategic relationship, or customer-trust judgment. |  
| `3` stale sales prospects `[Observed from brief]` | `yellow\_owner` | Pipeline decay has time value but usually does not need CEO review. | Check CRM/HubSpot/Salesforce stage, last reply, Gong/call context if available, account fit, relationship risk, owner, and trigger. Re-entry note or mark dead. | Only named account / Eric relationship. |  
| `23` pending review items `[Observed from brief]` | Batch through score gate. | Age is a symptom; contents decide priority. | Sort into queue statuses and pull only `3-5` red items `[Observed from brief]` for Eric. | Red items only. |  
| `15` Oracle quick-win keywords `[Observed from brief]` | `needs\_source\_check` -> `green\_auto` or `yellow\_owner` | SEO upside exists only if source, intent, existing page, and owner check out. | Verify GSC/GA4 data, existing-page fit, intent, effort, CMS/ClickFlow state, and convert winners into roadmap items. | Only strategic page / brand risk. |

## 5. Eric's 30-Minute Calibration Block

| Minute | Ask | Output |  
|---|---|---|  
| `0-5` `[Assumed]` | "I am handling churn, stale prospects, and queue routing first. Any reversal?" | Priority confirmation. |  
| `5-12` `[Assumed]` | "Which actions are never auto-approved?" | Red-line list. |  
| `12-18` `[Assumed]` | "What can I decide without you for `7` days `[Assumed]`?" | Temporary delegation boundary. |  
| `18-24` `[Assumed]` | Review only `3-5` `red\_eric` items `[Observed from brief]`. | Decisions / deferrals. |  
| `24-30` `[Assumed]` | "Which metric should be on tomorrow morning's briefing?" | CEO leverage metric. |

## 6. Score-To-Route Gate

Separate value from risk. A blended score should never hide an irreversible or reputation-sensitive item.

### Action Value Score

Each field is scored `0/1/2` `[Assumed]`; max value score is `14` `[Assumed]`.

| Field | 0 | 1 | 2 |  
|---|---|---|  
| Revenue proximity | No revenue tie | Indirect revenue / productivity | Active revenue, pipeline, churn, retention |  
| Time sensitivity | Can wait a week | Needs this week | Degrades within `24-48` hours `[Assumed]` |  
| Source confidence | Vague / unverified | Partial source | Source record available |  
| System writeback | No writeback path | Manual update required | CRM/CMS/queue writeback available |  
| Owner clarity | No owner | Owner category known | Named owner can act today |  
| Effort size | >1 day `[Assumed]` | Half-day `[Assumed]` | <1 hour `[Assumed]` |  
| Learning value | One-off | Some pattern value | Improves future routing/prompt/rule |

### Risk Flags

| Flag | Route pressure | Examples |  
|---|---|---|  
| Irreversible customer-facing commitment | Red | Pricing, refund, contract, public promise |  
| Strategic relationship / named account | Red | CEO relationship, enterprise logo, partner |  
| Brand/public reputation risk | Red | Strong POV, public figures, competitors |  
| Legal/privacy/security risk | Red | PII, data access, compliance |  
| High financial exposure | Red | Spend changes, custom deal structure |  
| Novel action with no precedent | Yellow | First live execution of a workflow |  
| Low source confidence | Yellow | Agent lacks source record |  
| Reversible routine action | Green possible | Internal brief, approved-template draft, owner reminder |

### Routing Rules

| Route | Criteria | Examples | Logging |  
|---|---|---|---|  
| Green | Value `>=7` `[Assumed]`, no red flags, source verified, reversible, known pattern. | Internal brief, approved-template content, routine SEO title/internal-link task. | Auto-execute; daily summary. |  
| Yellow owner | Value `>=5` `[Assumed]` or uncertainty without CEO-specific risk. | Churn investigation, new-page SEO recommendation, non-ICP prospect, sales re-entry note. | Owner review within `4` hours `[Assumed]`; outcome logged. |  
| Yellow council | High-value ambiguity where expert judgment improves the call. | Creative refresh, channel pivot, SEO/content positioning, paid media restructure. | Council memo; GM decision. |  
| Red Eric | Any red flag or CEO-specific context. | Pricing, strategic account, public controversy, irreversible customer promise, senior people issue. | `3-5` item Eric queue `[Observed from brief]`. |  
| Kill/defer | Value `<5` `[Assumed]` and no time-sensitive downside. | Daily competitor pricing noise, low-signal candidate follow-up. | Rejection reason logged. |

## 7. Agent-Specific Examples

| Agent | Green | Yellow | Red |  
|---|---|---|---|  
| Oracle | Existing-page metadata/internal-link quick win with verified GSC/GA4 data and CMS/ClickFlow path. | New page, mixed intent, material rewrite. | Public SEO/AEO claim, strategic positioning page, customer-sensitive case study. |  
| Flash | Repurpose approved video into template-matched post. | New hook pattern, strong POV, competitor mention. | Public figure, partner, legal/brand-sensitive claim. |  
| Cyborg | Follow-up inside approved sequence for non-strategic prospect with CRM writeback. | New ICP edge case, high-fit senior buyer, unusual personalization. | Eric relationship, named account, offer/pricing change. |  
| Alfred | Internal briefing, owner reminder, cross-agent status summary. | Churn investigation, deal revival, workflow proposal. | New core cron, killing core cron, pricing/client escalation. |

## 8. Cron Audit Framework

Audit question: does this cron create accepted work or review debt?

Score each cron by action rate, revenue proximity, review burden, source quality, owner clarity, false-positive cost, false-negative cost, and memory feedback "[Assumed]".

| Cron | Current signal | Verdict | Reason | 30-day change |  
|---|---|---|---|---|  
| GSC Quick Win Scan | "12%" actioned "[Observed from brief]" | Improve | Potentially valuable, but current volume creates noise. | Surface only verified existing-page, commercial-intent opportunities with clear owner and low effort. |  
| LinkedIn Prospect Sourcer | "8%" meetings "[Observed from brief]" | Improve | Pipeline-relevant, but "2x/day" "[Observed from brief]" likely floods review. | Tighten ICP/trigger/relationship score; batch daily "[Assumed]"; require CRM/HubSpot writeback and pipeline-impact read. |  
| Content Repurpose - YouTube | "35%" published "[Observed from brief]" | Keep / improve | Editorial rejection is healthy; needs winner/loser memory. | Feed accepted/rejected hooks into memory weekly "[Assumed]". |  
| Deal Revival Scan | "60%" followed up "[Observed from brief]" | Keep | Revenue-proximate, low-volume, clear owner action. | Require follow-up outcome logging. |  
| Competitor Pricing Monitor | "3%" action "[Observed from brief]" | Kill / pause | Daily pricing scrape creates rare action and executive distraction. | Replace with monthly or exception-based scan "[Assumed]". |  
| Churn Signal Detector | "20%" investigated "[Observed from brief]" | Improve urgently | False negatives are expensive; response loop is broken. | Same-day yellow escalation with owner, source links, customer state, and save-play status "[Assumed]". |  
| Weekly Content Calendar | "50%" used "[Observed from brief]" | Keep / improve | Useful planning artifact if tied to accepted work. | Track publish/outcome and feed winners back into memory. |  
| Candidate Outreach Follow-up | "5%" response "[Observed from brief]" | Kill / redesign | Low response plus brand risk from robotic follow-up. | Kill auto-follow-up; weekly top-"5" human review only "[Assumed]". |

## 9. Revenue/FTE Dashboard

This is the monthly GM view that keeps AI leverage tied to business health, not output volume.

| Metric | Source | Why it matters |  
|---|---|---|  
| Revenue per FTE | Finance + people data | Primary GM mandate from job post. |  
| TES revenue mix | Finance + ClickFlow/Karrot/product data | Tracks progress toward "50%" TES revenue by 2028 "[Observed from brief]". |  
| Net margin | P&L | Prevents artificial shrinkage and hidden service debt. |  
| Accepted agent work | Queue log | Separates production from noise. |  
| CEO red decisions/day | "red\_eric" queue | Measures founder leverage. |  
| Churn flags investigated same day | ClickFlow/product + queue | Protects TES revenue. |  
| Pipeline re-entry actions completed | CRM + queue | Recovers stale demand. |  
| Cost per accepted action | model/tool spend + queue | Prevents hidden AI/tool bloat. |  
| Writeback completion rate | CRM/CMS/queue audit | Prevents useful Slack messages from dying outside the system of record. |

## 10. Mistake Handling

When a green item should not have auto-executed:

1. Reverse if reversible.
2. Notify the owner in "3" lines "[Assumed]": what happened, current risk, correction.
3. Move the item type from green to yellow.
4. Add a memory record with failed rule, source, route, impact, and new threshold.
5. Review weekly "[Assumed]"; "2" repeated failures "[Assumed]" in one category means the rule is wrong.

## 11. Memory Record Schema

```
```yaml
item_id:
agent:
source_links:
system_of_record:
input_summary:
route: green_auto | yellow_owner | yellow_council | red_eric | killed
route_reason:
owner:
action_taken:
writeback_status:
audit_link:
outcome_metric:
cost_or_time_spend:
failure_or_exception:
rule_update:
next_review_date:
```
```

## 12. Operating Cadence

| Cadence | Artifact / meeting | Purpose |  
|---|---|---|  
| Daily "7:30 AM" "[Assumed]" | CEO leverage briefing | "3-5" red decisions "[Observed from brief]", churn/pipeline exceptions, yesterday's accepted work. |  
| Daily "4:30 PM" "[Assumed]" | Queue closeout | Zero unowned yellow/red items; all executions logged. |  
| Weekly "[Assumed]" | Cron audit / rule tuning | Kill, pause, tighten, or graduate crons based on outcomes. |  
| Weekly "[Assumed]" | Memory sweep | Convert repeated decisions into rules, prompts, SOPs, or kill criteria. |  
| Monthly "[Assumed]" | Revenue/FTE review | Tie AI leverage to revenue, retention, margin, and productivity. |  
| Monthly "[Assumed]" | Model/tool spend review | Check model hierarchy, API fallback use, and cost per accepted action. |

## 13. SSR Trust Check

Borrow the public Single Brain operating standard "[Observed from public websites]" and make it operational:

| Standard | GM test |  
|---|---|  
| Security | Correct data scopes, no client data bleed, credential hygiene, role-based access. |  
| Stability | Agents do not drift, repeat bad actions, or require constant rescue. |  
| Reliability | Owners see clear status, completed work, audit trail, and response times fast enough to trust. |

#### ## 14. What Stays Human

- Pricing, discounts, refunds, contract terms, and spend-risk decisions.
- Strategic account communication and client retention conversations where trust is at risk.
- Public claims involving customers, partners, competitors, public figures, or legal/regulatory issues.
- Hiring/firing, senior candidate persuasion, and people conflict.
- Automation rules that have not survived live review.
- Eric-specific taste, vision, partnerships, and strategic priority tradeoffs.

#### ## 15. What Breaks It

- Source data is wrong, stale, missing, or inaccessible.
- CRM stages, GSC/GA4 access, enrichment, CMS state, or account ownership are dirty before the agent acts.
- Agents post useful Slack messages but do not write outcomes back to CRM/CMS/queue.
- Agents optimize output volume instead of accepted work.
- Yellow items have no owner or SLA.
- Eric keeps reviewing non-red decisions.
- Councils become delay theater instead of high-value judgment.
- Memory records are not updated, so mistakes repeat.
- Churn/pipeline items are treated like ordinary queue items.
- Client/data isolation or permission scope is unclear.
- Model/tool spend is unmanaged, with no hierarchy or cost per accepted action.
- The dashboard becomes the work instead of a trigger for execution.

#### ## 16. Evidence Hooks

| Claim in packet | Proof source |

|---|---|

| I have moved revenue/FTE and margin together. | QuickBooks P&L snapshots and FTE model summarized in evidence log. |

| I use gated scoring / routing patterns. | x-accel gate/scoring source records and sanitized workflow summary. |

| I use expert councils for ambiguity. | Decision Council workflow and safe/anonymized council artifact. |

| I maintain durable operating memory. | SloaneVault/qmd index, lessons, project notes, and decision logs. |

| I have built sales and delivery automation. | HG Market Report and HG SEO Commander summaries. |

| I have operated a real scheduled-agent / cron fleet. | Hermes/open-agents scheduler manifests and live state summarized in the evidence log. |

| I understand maturity limits. | Explicit distinction among x-accel, SloaneVault, and early-stage HeavisideOS. |

# score-review

## # Ops COO 009 - Packet v2 Score Review

Status: score review after rebuilding the v2 packet, adding public-website calibration, adding the Eric/Chris operating-profile match, adding redacted proof summaries, completing the council remediation passes, adding explicit talent-leadership / accountability proof, adding the PASCO/full-time-role delegation proof, adding the less-than-`5`-hours/week founder-management claim, and adding line-by-line job-post/challenge-compliance scoring.

Last rescored: 2026-05-14.

Reviewed files:

- `tasks/final-upload-v2/written-answer.md`
- `tasks/final-upload-v2/written-answer.pdf`
- `tasks/final-upload-v2/operating-artifact.md`
- `tasks/final-upload-v2/evidence-log.md`
- `tasks/final-upload-v2/ai-usage-disclosure.md`
- `tasks/final-upload-v2/candidate-profile.md`
- `tasks/final-upload-v2/attachment-map.md`
- `tasks/final-upload-v2/job-post-line-scorecard.md`
- `tasks/final-upload-v2/challenge-compliance-audit.md`
- `tasks/final-upload-v2/application-answers.md`
- `tasks/final-upload-v2/cover-letter.md`
- `tasks/final-upload-v2/hermes-closed-loop-agent-proof.md`
- `tasks/final-upload-v2/heavyside-org-chart-redacted-summary.md`
- `tasks/final-upload-v2/heavyside-rpe-profitability-redacted-summary.md`
- `tasks/final-upload-v2/heavyside-sales-delivery-automation-redacted-summary.md`
- `tasks/final-upload-v2/heavysideos-slack-eval-redacted-summary.md`
- `tasks/final-upload-v2/hermes-cron-fleet-redacted-summary.md`

Scoring sources:

- Live GM job post: `https://careers.singlegrain.com/jobs/567472-general-manager-agency`, rechecked 2026-05-14.
- `challenges/ops-coo-009/claude\_baseline.md`
- `challenges/ops-coo-009/scoring\_rubric.md`
- `SCORING.md`

## ## Mechanical QA

| Check | Result |  
|---|---|  
| Written answer word count | `1724` words |  
| Written answer compact PDF | `tasks/final-upload-v2/written-answer.pdf` |  
| PDF page count | `2` pages |  
| QA helper | Pass |  
| Candidate/blocker language | Pass |  
| Required sections | Pass: What stays human, What breaks, AI usage, Evidence, MMPI, LinkedIn |  
| Local path scan | Pass |  
| Sensitive-info scan | Warning only: redaction warnings in evidence/proof files; no secrets or raw client data detected by helper. |

## ## Score Posture

| Rubric gate | Read |  
|---|---|  
| Real judgment | Strong. Clear triage order, hard cron kills, CEO decision boundary, and explicit tradeoffs. |  
| Execution artifact | Strong. Artifact includes command queue, gate, cron audit, dashboard, memory schema, cadence, human boundary, and failure modes. |  
| Source-labeled specificity | Strong. Written answer and artifact label substantive numbers. Final human review should still scan after edits. |  
| Proof quality | Strong. Evidence log now points to a Tier 4 RPE/profitability artifact, a Tier 4 HeavysideOS eval artifact, Tier 2-3 sales/delivery automation artifacts, Tier 3 Hermes scheduler-fleet evidence, and Tier 3 public-roster / ownership-map talent proof. |  
| AI fluency with limits | Strong. Shows automation, source checks, routing, memory, councils, and human boundaries without overclaiming HeavysideOS maturity. |  
| Failure handling | Strong. Written answer and artifact both name failure modes and response mechanisms. |  
| Communication | Strong. Written answer is direct and fits `2` pages. Artifact is longer but inspectable. |  
| Public-site fit | Strong. Recommendations now reflect Single Grain's public model: AI execution inside Slack/CRM/analytics/CMS, Single Brain command/fleet architecture, ClickFlow roadmap workflows, Karrot/HubSpot pipeline writeback, and SSR trust standards. |  
| CEO operating fit | Strong. Packet now addresses the brief's explicit Eric profile: high Quick Start, low Follow Thru, weak workflow consistency, `10` starts / `3` finishes, and the need for a GM who turns starts into closed loops. |  
| Talent leadership / accountability | Stronger. Written answer and candidate profile now state the founder-leverage proof directly: Chris sustained and grew the agency while holding a full-time PASCO role, and his existing businesses take less than `5` hours/week to manage because he hires competent operators, gives ownership, avoids micromanagement, and intervenes where ownership or performance requires it. The safer frame is delegated operating capacity, not divided attention. |  
| TES strategy | Stronger. Written answer now names the `50%` TES revenue by 2028 goal and connects revenue/FTE to productized/tech-enabled revenue, accepted agent work, and disciplined core services. |  
| Complexity reduction | Good but watch. The thesis and artifact reduce review debt, but the support stack can look complex. Use the lean upload hierarchy and keep Hermes/x-accel/councils as supporting proof unless requested. |

## ## GM Job Requirement Alignment

Live job-post facts rechecked 2026-05-14: Single Grain wants a high-bar GM who builds operational leverage, moves revenue per employee from `\$180K` to `\$300K-\$400K`, reduces CEO dependency, operates the agency, enforces accountability, manages profitability and hiring discipline, builds leadership teams, and creates standardized processes, KPI dashboards, performance-management systems, and AI-enabled workflows.

| Job-post requirement | Fit score | Read |  
|---|---|---|  
| Builder of operational leverage, not traditional agency operator | `9/10` | Strong. Packet leads with RPE/margin, operating systems, command queue, gates, memory, and automation-to-accepted-work. |  
| Increase revenue per employee while maintaining business health | `8.75/10` | Strong but not complete proof at Single Grain scale. Heavyside shows RPE and margin improvement together; target operating model is clear. Needs live SG economics in first 30 days. |  
| Reduce CEO operational dependency | `9.5/10` | One of the strongest areas. The whole artifact is designed around reducing Eric to red decisions and CEO-level judgment. |

| Operated or scaled a performance marketing agency / service business | `8.5/10` | Strong founder/operator proof since 2011, but smaller than Single Grain and partially applicant-stated. |  
 | Built operational systems that improved productivity | `9.25/10` | Strong. HG Market Report, HG SEO Commander, HeavisideOS eval, Hermes scheduler, and RPE/profitability proof reinforce each other. |  
 | Demonstrated financial discipline | `8.75/10` | Strong. P&L-based RPE/margin history is exactly job-relevant. Could be even stronger with hiring-decision examples and budget governance examples. |  
 | Built leadership teams capable of operating independently | `8.75/10` | Improved. Public roster, ownership map, org chart, PASCO/full-time-role proof, and less-than-`5`-hours/week management claim now support this. Remaining weakness is lack of third-party/operator references. |  
 | Talent leadership and accountability | `8.75/10` | Stronger than before. The hiring/delegation model is clear; still could use one concrete underperformance/coaching example. |  
 | Systems and AI infrastructure | `9.5/10` | Very strong. This is probably the clearest match to the role's "tech-enabled services" direction. |  
 | Ability to take the role without dropping existing businesses | `8.75/10` | Stronger with the less-than-`5`-hours/week management claim. Keep the emphasis on delegated operating capacity and explicit boundaries, not casual split attention. |

Estimated job-post fit after the latest changes: `8.85-9.15/10`.

Detailed line-by-line job-post scoring:

- `job-post-line-scorecard.md`

Detailed challenge-part and win/lose audit:

- `challenge-compliance-audit.md`

## ## Internal Score Read

Not an official point score, but the current packet posture against the public rubric:

| Rubric area | Current read |  
 |---|---|  
 | Real judgment | `9/10` |  
 | Execution artifact | `9/10` |  
 | Source-labeled specificity | `9/10` |  
 | Proof quality | `9/10` |  
 | AI fluency with limits | `9-9.25/10` |  
 | Failure handling | `8.75-9/10` |  
 | Communication | `8.25-8.75/10` |  
 | Job-post fit | `8.85-9.15/10` |  
 | Baseline-beating likelihood | High, assuming reviewer accepts the redacted evidence summaries and attachments as credible and the existing-businesses point is framed as delegated operating capacity. |

## ## Second Council Packaging Review

The second council did not recommend a rewrite. It converged on a packaging decision:

- Freeze the written answer except typo, sensitivity, or label fixes.
- Do not upload the full proof stack by default.
- Lead with the role-relevant proof: written answer, operating artifact, evidence log, RPE/profitability proof, sales/delivery automation proof, AI disclosure, and candidate profile.
- Treat HeavisideOS and Hermes as appendix-only proof if the portal supports a clean appendix; otherwise keep them for walkthrough or interview follow-up.
- Keep x-accel, decision council, and SloaneVault raw proof as backup unless requested.

The main remaining risk is not substance. It is attachment bloat: too many internal tool names can make strong evidence feel like an internal catalog.

## ## Proof Artifact Upgrade

The new attachment set improves proof quality because it separates four different proof jobs:

| Proof job | Artifact | Scoring effect |  
 |---|---|---|  
 | GM economic proof | `heaviside-rpe-profitability-redacted-summary.md` | Moves the role-level leverage claim into Tier 4 before/after data. |  
 | Measured AI workflow proof | `heavisideos-slack-eval-redacted-summary.md` | Shows benchmark, remediation, gate improvement, and retrieval lift with real eval counts. |  
 | Built-system proof | `heaviside-sales-delivery-automation-redacted-summary.md` | Shows real sales/delivery systems and source-record basis; time-savings claims remain estimate-labeled. |  
 | Cron-fleet operating proof | `hermes-cron-fleet-redacted-summary.md` | Shows Chris operates a monitored `88`-job Hermes scheduler surface with routing, status, pause/error visibility, and fleet health checks; use as bounded Tier 3 proof, not as proof every cron is valuable. |

This materially reduces the risk that the packet reads like a polished strategy doc. It now has role-level economic proof, AI workflow proof, software-building proof, and bounded cron-fleet operating proof.

## ## Baseline Comparison

Claude is competent on triage, risk tiering, and cron audit. The v2 packet should beat it because it adds:

- real revenue/FTE and margin proof,
- real sales and delivery automation examples,
- redacted proof artifacts backing the strongest claims,
- a more inspectable artifact,
- a source-labeled evidence log,
- a real self-assessment,
- explicit founder-leverage and revenue/FTE framing tied to the GM posting.
- public-source specificity from Single Grain / Single Brain / ClickFlow / Karrot instead of generic agency-ops recommendations.
- explicit wiring complement: Chris's INTP / `6w5`, strong Type `1`, Kolbe `9-4-3-2` profile as the follow-through/source-checking counterweight to Eric's ENTJ / `8w7`, Quick Start `9`, Follow Thru `2` pattern.
- bounded cron-fleet operating evidence from Hermes/open-agents rather than only a proposed cron-audit framework.
- sharper operating-experience wording around replacing manual agency drag with source-data -> generated deliverable -> operator review -> customer send/publish -> outcome logging.
- explicit talent-leadership and founder-leverage wording: agency sustained and grew while Chris held a full-time PASCO role, and current businesses take less than `5` hours/week to manage, backed by hiring competent operators, ownership, cadence, and selective intervention.

## Claude Baseline Gap

| Challenge area | Claude baseline | Current packet |  
|---|---|---|  
| First 48 hours | Competent triage and stakeholder plan. | Similar triage quality, but stronger CEO-load design, explicit queue states, source labels, and no-unowned-item standard. |  
| Risk tiering | Good green/yellow/red categories and examples. | Stronger because value/risk are separated, red flags override value, and failures update memory/rules. |  
| Cron audit | Good sample judgments. | Stronger because it audits accepted work, review debt, source quality, owner clarity, and ties to Chris's own Hermes scheduler-fleet experience. |  
| Operating experience | Claude explicitly lacks real examples. | Strong human advantage: Heaviside, HG Market Report, HG SEO Commander, adoption example, RPE/margin history, PASCO parallel operating proof. |  
| MMP/ self-assessment | Claude cannot do this. | Strong human advantage: actual Kolbe/Enneagram/MBTI profile tied to Eric's operating gaps. |  
| Evidence/proof | Claude has no proof artifacts. | Strong advantage: evidence log, RPE/margin summary, system summaries, eval proof, org chart, AI disclosure. |

Current read: this should beat Claude decisively on Part 4 and proof quality, and at least match or beat it on Parts 1-3.

## Remaining Non-Scoring Blockers

- 1. Applicant should confirm the LinkedIn URL in `tasks/final-upload-v2/candidate-profile.md`.
- 2. Applicant should approve the self-assessment wording.
- 3. Applicant should approve the lean upload set versus appendix set in `tasks/final-upload-v2/attachment-map.md`.
- 4. Applicant should approve the existing-businesses framing: delegated operating capacity and explicit boundaries, not divided attention.
- 5. Re-run live posting check immediately before actual upload if time passes.

## Verdict

The v2 packet is now materially stronger than Claude's published baseline and aligned with the public scoring guide. It is not merely a better essay; it has an inspectable operating artifact, a proof-tiered evidence log, and real operating metrics.

# SUBMISSION-COMPLIANCE-README

---

# Ops COO 009 - Submission Compliance README

Status checked: 2026-05-14.

## ## Sources Checked

- Challenge brief: <https://github.com/ericosiu/beat-claude/blob/main/challenges/ops-coo-009/brief.md>.
- Public scoring guide: <https://github.com/ericosiu/beat-claude/blob/main/SCORING.md>.
- Challenge-specific rubric: [https://github.com/ericosiu/beat-claude/blob/main/challenges/ops-coo-009/scoring\\_rubric.md](https://github.com/ericosiu/beat-claude/blob/main/challenges/ops-coo-009/scoring_rubric.md).
- Claude baseline: [https://github.com/ericosiu/beat-claude/blob/main/challenges/ops-coo-009/claude\\_baseline.md](https://github.com/ericosiu/beat-claude/blob/main/challenges/ops-coo-009/claude_baseline.md).
- Live job post: <https://careers.singlegrain.com/jobs/567472-general-manager-agency>.

## ## Page-Limit Interpretation

The job post says to submit a PDF or Markdown response, max 4 pages.

The challenge brief itself is stricter: maximum 2 pages. It also says diagrams and decision trees do not count toward the limit.

The public scoring guide says most challenges have a 4-page written-answer cap. It also says diagrams, artifact links, code, sheets, and short Loom demos do not count unless the brief says otherwise.

Safe interpretation: treat `written-answer.pdf` as the official written answer and keep it to 2 pages.

Treat operating artifacts, evidence logs, source labels, disclosures, proof files, screenshots/images, and appendices as supporting materials. Do not treat them as extra written-answer pages.

## ## What To Submit

Best upload path if the portal allows multiple files:

1. `00\_OFFICIAL\_challenge\_answer\_2\_pages.pdf`
2. `operating-artifact.md`
3. `evidence-log.md`
4. `ai-usage-disclosure.md`
5. `candidate-profile.md`
6. Core proof summaries, if the portal accepts supporting files

If the portal only accepts one challenge file, the safer file is still `00\_OFFICIAL\_challenge\_answer\_2\_pages.pdf`. Use the cover letter, public page, or application fields to point reviewers to the supporting materials.

`01\_FULL\_REVIEW\_packet\_with\_image\_appendix.pdf` is a review/backup packet. It contains the official 2-page written answer plus image-rendered supporting artifacts.

Use it only if you are comfortable with the appendix risk. The risk: the portal or reviewer might count every PDF page. The brief and scoring guide allow supporting material outside the written limit, but the upload UI may not make that obvious.

## ## Required Packet Coverage

- Written answer: `00\_OFFICIAL\_challenge\_answer\_2\_pages.pdf` and `written-answer.md`.
- Operating artifact: `operating-artifact.md`.
- Evidence log: `evidence-log.md`.
- Number source labels: included in `written-answer.md`, `operating-artifact.md`, and `evidence-log.md`.
- AI usage disclosure: `ai-usage-disclosure.md`.
- What breaks it: included in `written-answer.md` and `operating-artifact.md`.
- What stays human: included in `written-answer.md`, `operating-artifact.md`, and `ai-usage-disclosure.md`.
- LinkedIn URL: `candidate-profile.md`.
- MMPI-style self-assessment: `candidate-profile.md` and `written-answer.md`.

## ## Do Not Upload By Default

- Raw QuickBooks exports.
- Raw client data.
- Raw internal logs.
- Internal-only score reviews unless asked.

The included proof files are redacted summaries intended to support a live walkthrough without exposing unnecessary confidential detail.